

# SciencesPo Computational Economics

## Spring 2019

Florian Oswald

April 15, 2019

### 1 The Julia-is-Fast Benchmark Fun: sum

(This material began life as a wonderful lecture by Steven Johnson at MIT: <https://github.com/stevengj/18S096/blob/master/lectures/lecture1/Boxes-and-registers.ipynb>.)

In this notebook we are going to compare performance of a very simple function across several different languages: The sum. This function computes

$$\text{sum}(a) = \sum_{i=1}^n a_i$$

where  $n$  is the length of  $a$ . Let's get a vector of numbers:

```
In [1]: a = rand(10^7) # 1D vector of random numbers, uniform on [0,1)
```

```
Out[1]: 10000000-element Array{Float64,1}:
```

```
0.1663319579196314
0.30895312060879365
0.6179474041044644
0.06127644374834973
0.2573021832089899
0.16796863745240764
0.27852662730874767
0.5086510453252497
0.1560114763009781
0.7288501465239252
0.11324653454998068
0.12629501636416296
0.9000757666639296

0.7236680378178086
0.8672746082443499
0.6332323213343003
0.4469263836490769
0.6862137821151018
0.9506520059092591
```

```

0.17982818693745917
0.06950168957092684
0.7418679573560714
0.6435725623459287
0.9592755512624958
0.9511577329348149

```

We would expect to see  $0.5 \cdot 10^7$ , since each element has an expected value of 0.5.

## Benchmarks in different languages

We will use BenchmarkTools.jl for this exercise. This is because the standard @time macro suffers from sample bias:

```
In [2]: @time sum(a)
```

```
@time sum(a)
```

```
@time sum(a)
```

```

0.045626 seconds (94.66 k allocations: 4.730 MiB)
0.004711 seconds (5 allocations: 176 bytes)
0.004619 seconds (5 allocations: 176 bytes)

```

```
Out[2]: 5.000250466313629e6
```

```
In [3]: using BenchmarkTools
```

## 1.1 C is what you have to beat

C is often considered the gold standard: difficult on the human, nice for the machine. Getting within a factor of 2 of C is often satisfying. Nonetheless, even within C, there are many kinds of optimizations possible that a naive C writer may or may not get the advantage of.

One can compile C code in Julia. Note that the `"""` wrap a multi-line string.

```

In [4]: C_code = """
#include <stddef.h>
double c_sum(size_t n, double *X) {
    double s = 0.0;
    for (size_t i = 0; i < n; ++i) {
        s += X[i];
    }
    return s;
}
"""

const Clib = tempname() # make a temporary file

# compile to a shared library by piping C_code to gcc

```

```

# (works only if you have gcc installed):
using Libdl

open(`gcc -fPIC -O3 -ffast-math -msse3 -xc -shared -o $(Clib * "." * Libdl.dlext) -`,
      print(f, C_code)
end

# define a Julia function that calls the C function:
c_sum(X::Array{Float64}) = ccall(("c_sum", Clib), Float64, (Csize_t, Ptr{Float64}), length(X), X)

c_sum(a)

c_sum(a) sum(a) # type \approx and then <TAB> to get the symbol

c_bench = @benchmark c_sum($a)

println("C: Fastest time was $(minimum(c_bench.times) / 1e6) msec")

d = Dict{String,Float64}() # a "dictionary", i.e. an associative array
d["C"] = minimum(c_bench.times) / 1e6 # in milliseconds
d

```

C: Fastest time was 4.414939 msec

```

Out[4]: Dict{Any,Any} with 1 entry:
  "C" => 4.41494

```

We can see above that the BenchmarkTools library takes many sample runs to account for machine noise in the benchmark. We can look at the distribution of times:

```

In [5]: using Plots

```

```

t = c_bench.times / 1e6 # times in milliseconds
using Statistics
m, = minimum(t), std(t)


histogram(t, bins=500,
           xlim=(m - 0.01, m + 0.01),
           xlabel="milliseconds", ylabel="count", label="")

```

```

Out[5]:

```



## 1.2 Next: python's built-in sum

In [6]: `using PyCall`

```
# Call a low-level PyCall function to get a Python list, because
# by default PyCall will convert to a NumPy array instead (we benchmark NumPy below):
```

```
apy_list = PyCall.array2py(a, 1, 1)
```

```
# get the Python built-in "sum" function:
pysum = pybuiltin("sum")
```

```
pysum(a)
```

```
pysum(a)    sum(a)
```

```
py_list_bench = @benchmark $pysum($apy_list)
```

```
d["Python built-in"] = minimum(py_list_bench.times) / 1e6
d
```

```
Out[6]: Dict{Any,Any} with 2 entries:
  "C" => 4.41494
  "Python built-in" => 54.7549
```

### 1.3 Next: python's numpy sum

Takes advantage of hardware “SIMD”, but only works when it works.

numpy is an optimized C library, callable from Python. It may be used within Julia as follows:

In [7]: `using` Conda

```
numpy_sum = pyimport("numpy")["sum"]
apy_numpy = PyObject(a) # converts to a numpy array by default

py_numpy_bench = @benchmark $numpy_sum($apy_numpy)

numpy_sum(apy_list) # python thing

numpy_sum(apy_list)    sum(a)

d["Python numpy"] = minimum(py_numpy_bench.times) / 1e6
d
```

```
Out [7]: Dict{Any,Any} with 3 entries:
  "C"          => 4.41494
  "Python numpy" => 4.6328
  "Python built-in" => 54.7549
```

### 1.4 Next: python hand-written

We could try and see how our hand-written implementation performs:

```
In [8]: py"""
def py_sum(a):
    s = 0.0
    for x in a:
        s = s + x
    return s
"""

sum_py = py"py_sum"

py_hand = @benchmark $sum_py($apy_list)

sum_py(apy_list)

sum_py(apy_list)    sum(a)

d["Python hand-written"] = minimum(py_hand.times) / 1e6
d
```

```
Out [8]: Dict{Any,Any} with 4 entries:
  "C"          => 4.41494
```

```

"Python numpy"          => 4.6328
"Python hand-written"   => 284.278
"Python built-in"       => 54.7549

```

## 1.5 julia built-in

julias library is written entirely in julia! No C at all! you can easily look at the code by typing

```
In [9]: @which sum(a)
```

```
Out[9]: sum(a::AbstractArray) in Base at reducedim.jl:645
```

```
In [10]: j_bench = @benchmark sum($a)
```

```

d["Julia built-in"] = minimum(j_bench.times) / 1e6
d

```

```
Out[10]: Dict{Any,Any} with 5 entries:
```

```

"C"          => 4.41494
"Python numpy"      => 4.6328
"Python hand-written" => 284.278
"Python built-in"   => 54.7549
"Julia built-in"    => 4.48438

```

## 1.6 julia hand-written

```
In [11]: function mysum(A)
```

```

    s = 0.0 # s = zero(eltype(A))
    for a in A
        s += a
    end
    s
end

```

```
j_bench_hand = @benchmark mysum($a)
```

```

d["Julia hand-written"] = minimum(j_bench_hand.times) / 1e6
d

```

```
Out[11]: Dict{Any,Any} with 6 entries:
```

```

"C"          => 4.41494
"Python numpy"      => 4.6328
"Julia hand-written" => 9.74541
"Python hand-written" => 284.278
"Python built-in"   => 54.7549
"Julia built-in"    => 4.48438

```

## 1.7 R built-in

```
In [12]: using RCall
```

```
    r_bench = @benchmark R"sum($a)"  
    d["R built-in"] = minimum(r_bench.times) / 1e6
```

```
Info: Recompiling stale cache file /Users/florian.oswald/.julia/compiled/v1.0/RCall/8GFyb.ji :  
@ Base loading.jl:1190
```

```
Out[12]: 48.683041
```

## 1.8 Summary

```
In [13]: for (key, value) in sort(collect(d), by=x->x[2])  
          println(rpad(key, 20, "."), lpad(round(value,digits= 2), 10, "."))  
        end
```

```
C...4.41  
Julia built-in...4.48  
Python numpy...4.63  
Julia hand-written...9.75  
R built-in...48.68  
Python built-in...54.75  
Python hand-written...284.28
```

### 1.8.1 Take aways (on my computer!):

1. C is fastest
2. built-in julia checks out very close to C, and ex-equo with the numpy
3. Hand written julia gets compiled to very efficient machine code in this example.
4. Python and R built-in sums are roughly 10 times slower than C, julia and numpy
5. Hand writing python code without any optimizations performs poorly in this instance.

```
In [ ]:
```