

optimization

March 24, 2017

1 Computational Economics: Unconstrained Optimization

Florian Oswald Sciences Po, 2017

1.1 Some Taxonomy and Initial Examples

- In most of the examples to follow, we talk about *minimization* of a function f . Everything we do also applies to maximization, since $\min_x f(x) = \max_x -f(x)$.
- Here is a generic optimization problem:

$$\min_{x \in \mathbb{R}^n} f(x) \text{ s.t. } \begin{array}{ll} c_i(x) = 0, & i \in E \\ c_i(x) \geq 0, & i \in I \end{array}$$

- This is a general way of writing an optimization problem. E are all indices as equality constraints, I are all inequality constraints.
- An example of such a problem might be

$$\min (x_1 - 2)^2 + (x_2 - 1)^2 \text{ s.t. } \begin{array}{l} x_1^2 - x_2 \leq 0 \\ x_1 + x_2 \leq 2 \end{array}$$

- Here is a picture of that problem taken from the textbook [nocedal-wright] (for copyright reasons, I cannot show this in the online version of the slides.):

1.2 Kinds of problems considered

- Don't talk about stochastic optimization methods:
 - Simulated Annealing
 - MCMC
 - other Stochastic Search Methods
 - A gentle introduction is [casella-R]




Figure 1.1 in [nocedal-wright]

1.3 Transportation Problem

A chemical company has two factories F_1, F_2 and a dozen retail outlets R_1, \dots, R_{12} . Each factory i can produce at most a_i tons of output each week. Each retail outlet j has a weekly demand of b_j tons per week. The cost of shipping from F_i to R_j is given by c_{ij} . How much of the product to ship from each factory to each outlet, minimize cost, and satisfy all constraints? let's call x_{ij} the number of tons shipped from i to j .

- A mathematical formulation of this problem is

$$\begin{aligned} & \min \sum_{ij} c_{ij} x_{ij} \\ & \text{subject to } \sum_{j=1}^{12} x_{ij} \leq a_i, \quad i = 1, 2 \\ & \sum_{i=1}^2 x_{ij} \geq b_j, \quad j = 1, \dots, 12 \\ & x_{ij} \geq 0, \quad i = 1, 2, j = 1, \dots, 12 \end{aligned}$$

- This is called a *linear programming* problem, because both objective function and all constraints are linear.
- With any of those being nonlinear, we would call this a non-linear problem.

1.4 Constrained vs Unconstrained

- There are many applications of both in economics.




Figure 1.2 in [nocedal-wright]

- Unconstrained: maximum likelihood
- Constrained: MPEC
- It is sometimes possible to transform a constrained problem into an unconstrained one.

1.5 Convexity

- Convex problems are easier to solve.
- What is convex?

A set $S \in \mathbb{R}^n$ is convex if the straight line segment connecting any two points in S lies entirely inside S . A function f is a convex function, if its domain S is a convex set, and for any two points $x, y \in S$, we have that

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

for all $\alpha \in [0, 1]$

- Simple instances of convex sets are the unit ball $\{y \in \mathbb{R}^n, \|y\|_2 \leq 1\}$, and any set defined by linear equalities and inequalities.
- *convex Programming* describes a special case of the introductory minimization problem where
 - the objective function is convex,
 - the equality constraints are linear, and
 - the inequality constraints are concave.

1.6 Optimization Algorithms

- All of the algorithms we are going to see employ some kind of *iterative* procedure.
- They try to improve the value of the objective function over successive steps.

- The way the algorithm goes about generating the next step is what distinguishes algorithms from one another.
 - Some algos only use the objective function
 - Some use both objective and gradients
 - Some add the Hessian
 - and many variants more

1.7 Desirable Features of any Algorithm

- Robustness: We want good performance on a wide variety of problems in their class, and starting from *all* reasonable starting points.
- Efficiency: They should be fast and not use an excessive amount of memory.
- Accuracy: They should identify the solution with high precision.

1.8 A Word of Caution

- You should **not** normally attempt to write a numerical optimizer for yourself.
- Entire generations of Applied Mathematicians and other numerical pro's have worked on those topics before you, so you should use their work.
 - Any optimizer you could come up with is probably going to perform below par, and be highly likely to contain mistakes.
 - Don't reinvent the wheel.
- That said, it's very important that we understand some basics about the main algorithms, because your task is **to choose from the wide array of available ones**.

2 Unconstrained Optimization: What is a solution?

- A typical unconstrained optimization problem will look something like this:

$$\min_x f(x), \quad x \in \mathbb{R}^n$$

and where $f : \mathbb{R}^n \mapsto \mathbb{R}$ is a smooth function.


- In general, we would always like to find a *global* minimizer, i.e. a point

$$x^* \text{ where } f(x^*) \leq f(x) \quad \forall x$$

- Since our algorithm is not going to visit many points in \mathbb{R}^n (or so we hope), we can never be totally sure that we find a global optimizer.
- Most optimizers can only find a *local* minimizer. That is a point

$$x^* \text{ where } f(x^*) \leq f(x) \quad \forall x \in \mathcal{N}$$

where \mathcal{N} is a neighborhood around x^* .



Global min at $f(512, 404.2319)$. By Gaortizg [GFDL](#) or [CC BY-SA 3.0](#), via Wikimedia Commons

2.1 Global minization can be very hard sometimes.

2.2 (Unconstrained) Optimization in Julia

- Umbrella Organisation: <http://www.juliaopt.org>
 - We will make ample use of this when we talk about constrained optimisation.
 - The Julia Interface to the very well established C-Library [NLOpt](#) is called [NLOpt.jl](#). One could use [NLOpt](#) without constraints in an unconstrained problem.
- [Roots.jl](#): Simple algorithms that find the zeros of a univariate function.
- Baseline Collection of unconstrained optimization algorithms: [Optim.jl](#)

2.3 Introducing [Optim.jl](#)

- Multipurpose unconstrained optimization package
 - provides 8 different algorithms with/without derivatives
 - univariate optimization without derivatives

2.4 The Golden Ratio or Bracketing Search for 1D problems

- A derivative-free method
- a Bracketing method
 - find the local minimum of f on $[a, b]$
 - select 2 interior points c, d such that $a < c < d < b$

- * $f(c) \leq f(d) \implies$ min must lie in $[a, d]$. replace b with d , start again with $[a, d]$
- * $f(c) > f(d) \implies$ min must lie in $[c, b]$. replace a with c , start again with $[c, b]$
- how to choose b, d though?
- we want the length of the interval to be independent of whether we replace upper or lower bound
- we want to reuse the non-replaced point from the previous iteration.
- these imply the golden rule:
- new point $x_i = a + \alpha_i(b - a)$, where $\alpha_1 = \frac{3-\sqrt{5}}{2}, \alpha_2 = \frac{\sqrt{5}-1}{2}$
- α_2 is known as the *golden ratio*, well known for it's role in renaissance art.

2.4.1 Bracketing Search in Julia

- The package `Optim.jl` provides an implementation of "Brent's Method" as well as the golden section search:

```
In [9]: using Plots
        using Optim
        f(x) = exp(x) - x^4
        minf(x) = -f(x)
        brent = optimize(minf, 0, 2, Brent())
        golden = optimize(minf, 0, 2, GoldenSection())
        println("brent = $brent")
        println("golden = $golden")
        plot(f, 0, 2)
```

```
brent = Results of Optimization Algorithm
* Algorithm: Brent's Method
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 12
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 13
```

```
WARNING: Method definition f{Any} in module Main at In[1]:2 overwritten at In[9]:3.
WARNING: Method definition minf{Any} in module Main at In[2]:3 overwritten at In[9]:4.
```

```
golden = Results of Optimization Algorithm
* Algorithm: Golden Section Search
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 37
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 38
```

```
In [2]: # how well does this do with many local minima?
fun(x) = exp(x) - x^4 + sin(40*x)
minf(x) = -fun(x)
grid = collect(0:0.0001:2);
v,idx = findmax(Float64[fun(x) for x in grid])
println("grid maximizer is $(grid[idx])")
golden = optimize(minf,0,2,GoldenSection())
brent = optimize(minf,0,2,Brent())
using Base.Test
println("brent minimizer = $(brent.minimizer)")
println("golden minimizer = $(golden.minimizer)")
plot(fun,0,2)
```

grid maximizer is 0.8247

WARNING: Method definition minf(Any) in module Main at In[1]:3 overwritten at In[2]:3.

UndefVarError: optimize not defined

2.5 Beyond One Dimension

2.5.1 Introducing Rosenbrock's Banana function

The Banana function is defined by

$$f(x, y) = (a - x)^2 + b(y - x^2)^2$$

2.5.2 What is the minimum of that function?

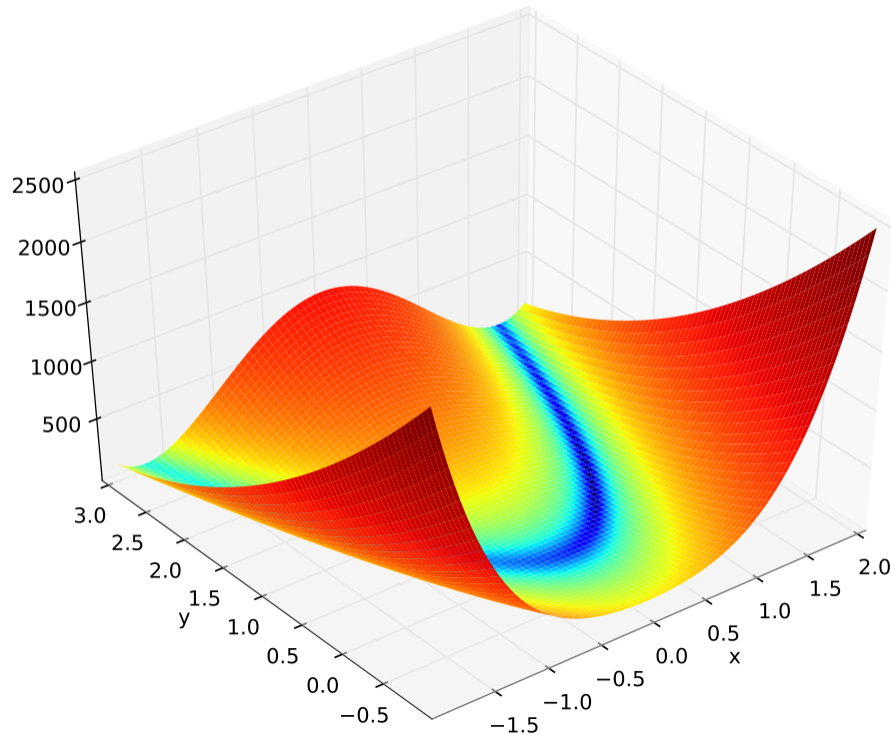
- For $a = 1, b = 100$, what is the global minimum of that function?
- What are the inputs one needs to supply to an algorithm in a more general example?

2.6 Rosenbrock Banana and Optim.jl

- We will use Optim for the rest of this lecture.
- We need to supply the objective function and - depending on the solution algorithm - the gradient and hessian as well.

```
In [3]: using Optim
        rosenbrock = Optim.UnconstrainedProblems.examples["Rosenbrock"]

        # contains:
        # function rosenbrock(x::Vector)
```



Banana for $a = 0$. By Gaortizg [GFDL](#) or [CC BY-SA 3.0](#), via Wikimedia Commons

```
#      return (1.0 - x[1])^2 + 100.0 * (x[2] - x[1]^2)^2
# end

# function rosenbrock_gradient!(x::Vector, storage::Vector)
#     storage[1] = -2.0 * (1.0 - x[1]) - 400.0 * (x[2] - x[1]^2) * x[1]
#     storage[2] = 200.0 * (x[2] - x[1]^2)
# end

# function rosenbrock_hessian!(x::Vector, storage::Matrix)
#     storage[1, 1] = 2.0 - 400.0 * x[2] + 1200.0 * x[1]^2
#     storage[1, 2] = -400.0 * x[1]
#     storage[2, 1] = -400.0 * x[1]
#     storage[2, 2] = 200.0
# end

# there are many other examples on Optim.UnconstrainedProblems
```

Out [3]: Optim.UnconstrainedProblems.OptimizationProblem("Rosenbrock",Optim.UnconstrainedProblems

2.7 Comparison Methods

- We will now look at a first class of algorithms, which are very simple, but sometimes a good starting point.
- They just *compare* function values.

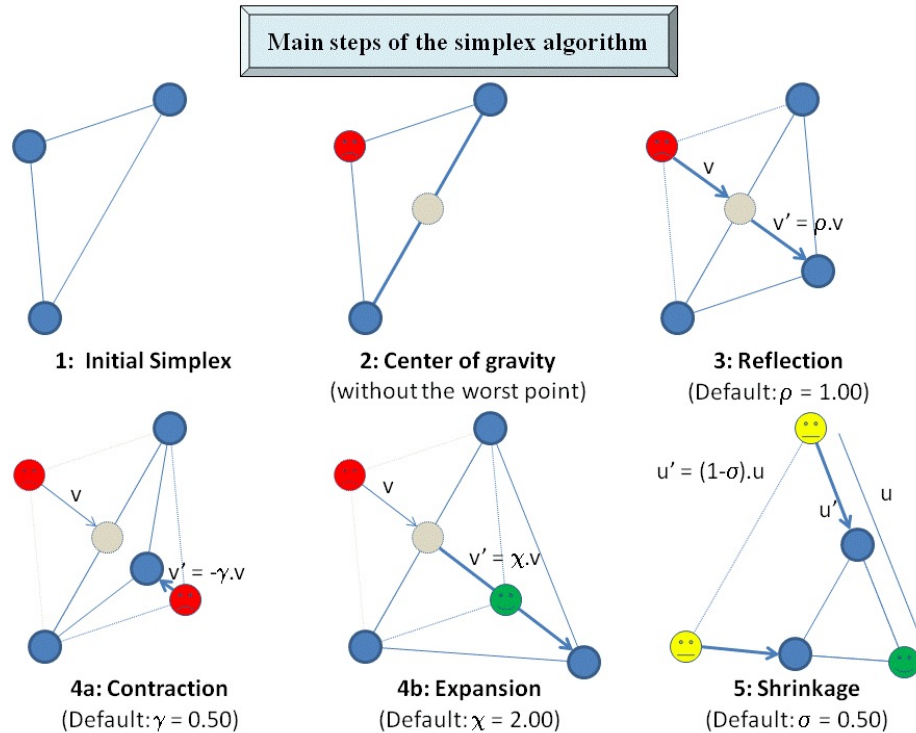
- *Grid Search* : Compute the objective function at $G = \{x_1, \dots, x_N\}$ and pick the highest value of f .
 - This is very slow.
 - It requires large N .
 - But it's robust (will find global optimizer for large enough N)

```
In [4]: # grid search on rosenbrock
        grid = collect(-1.0:0.1:3);
        grid2D = [[i;j] for i in grid,j in grid];
        val2D = map(rosenbrock.f,grid2D);
        r = findmin(val2D);
        println("grid search results in minimizer = $(grid2D[r[2]])")
```

```
grid search results in minimizer = [1.0,1.0]
```

2.8 Bracketing for Multidimensional Problems: Nelder-Mead

- The Goal here is to find the simplex containing the local minimizer x^*
- In the case where f is n -D, this simplex has $n + 1$ vertices
- In the case where f is 2-D, this simplex has $2 + 1$ vertices, i.e. it's a triangle.
- The method proceeds by evaluating the function at all $n + 1$ vertices, and by replacing the worst function value with a new guess.
- this can be achieved by a sequence of moves:
 - reflect
 - expand
 - contract
 - shrink movements.



- this is a very popular method. The matlab functions `fmincon` and `fminsearch` implements it.
- When it works, it works quite fast.
- No derivatives required.

In [5]: `optimize(rosenbrock, [0.0, 0.0], NelderMead())`

Out [5]: Results of Optimization Algorithm

```
* Algorithm: Nelder-Mead
* Starting Point: [0.0,0.0]
* Minimizer: [0.9999710322210338,0.9999438685860869]
* Minimum: 1.164323e-09
* Iterations: 74
* Convergence: true
* ((y-ŝ)/n < 1.0e-08: true
* Reached Maximum Number of Iterations: false
* Objective Function Calls: 108
```

- But.

2.9 Bracketing for Multidimensional Problems: Comment on Nelder-Mead

Lagarias et al. (SIOPT, 1999): At present there is no function in any dimension greater than one, for which the original Nelder-Mead algorithm has been proved to converge to a minimizer.

Given all the known inefficiencies and failures of the Nelder-Mead algorithm [...], one might wonder why it is used at all, let alone why it is so extraordinarily popular.