

SciencesPo Computational Economics Spring 2017

Florian Oswald

March 15, 2018

1 Optimization 2: Algorithms and Constraints

Florian Oswald Sciences Po, 2018

1.1 Bracketing

- A derivative-free method for *univariate* f
- works only on **unimodal** f
- (Draw choosing initial points and where to move next)

1.2 The Golden Ratio or Bracketing Search for 1D problems

- A derivative-free method
- a Bracketing method
 - find the local minimum of f on $[a, b]$
 - select 2 interior points c, d such that $a < c < d < b$
 - * $f(c) \leq f(d) \implies$ min must lie in $[a, d]$. replace b with d , start again with $[a, d]$
 - * $f(c) > f(d) \implies$ min must lie in $[c, b]$. replace a with c , start again with $[c, b]$
 - how to choose b, d though?
 - we want the length of the interval to be independent of whether we replace upper or lower bound
 - we want to reuse the non-replaced point from the previous iteration.
 - these imply the golden rule:
 - new point $x_i = a + \alpha_i(b - a)$, where $\alpha_1 = \frac{3-\sqrt{5}}{2}, \alpha_2 = \frac{\sqrt{5}-1}{2}$
 - α_2 is known as the *golden ratio*, well known for it's role in renaissance art.

```
In [1]: using Plots
        using Optim
        plotlyjs()
        f(x) = exp(x) - x^4
        minf(x) = -f(x)
        brent = optimize(minf, 0, 2, Brent())
        golden = optimize(minf, 0, 2, GoldenSection())
        plot(f, 0, 2)
```

```
gui()
println("brent = $brent")
println("golden = $golden")
```

WARNING: Method definition midpoints(Base.Range{T} where T) in module Base at deprecated.jl:56 o

WARNING: Method definition midpoints(AbstractArray{T, 1} where T) in module Base at deprecated.j

brent = Results of Optimization Algorithm

```
* Algorithm: Brent's Method
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 12
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 13
```

golden = Results of Optimization Algorithm

```
* Algorithm: Golden Section Search
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 37
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 38
```

1.2.1 Bisection Methods

- Root finding

1.3 Rosenbrock Banana and Optim.jl

- We can supply the objective function and - depending on the solution algorithm - the gradient and hessian as well.

In [2]: `using Optim`

```
rosenbrock = Optim.UnconstrainedProblems.examples["Rosenbrock"]
```

```
# contains:
# function rosenbrock(x::Vector)
#     return (1.0 - x[1])^2 + 100.0 * (x[2] - x[1]^2)^2
# end

# function rosenbrock_gradient!(x::Vector, storage::Vector)
#     storage[1] = -2.0 * (1.0 - x[1]) - 400.0 * (x[2] - x[1]^2) * x[1]
#     storage[2] = 200.0 * (x[2] - x[1]^2)
# end

# function rosenbrock_hessian!(x::Vector, storage::Matrix)
```

```

#     storage[1, 1] = 2.0 - 400.0 * x[2] + 1200.0 * x[1]^2
#     storage[1, 2] = -400.0 * x[1]
#     storage[2, 1] = -400.0 * x[1]
#     storage[2, 2] = 200.0
# end

# there are many other examples on Optim.UnconstrainedProblems

```

UndefVarError: UnconstrainedProblems not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

1.4 Comparison Methods

- We will now look at a first class of algorithms, which are very simple, but sometimes a good starting point.
- They just *compare* function values.
- *Grid Search*: Compute the objective function at $G = \{x_1, \dots, x_N\}$ and pick the highest value of f .
 - This is very slow.
 - It requires large N .
 - But it's robust (will find global optimizer for large enough N)

```

In [3]: # grid search on rosenbrock
grid = collect(-1.0:0.1:3);
grid2D = [[i;j] for i in grid,j in grid];
val2D = map(rosenbrock.f,grid2D);
r = findmin(val2D);
println("grid search results in minimizer = $(grid2D[r[2]])")

```

UndefVarError: rosenbrock not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

1.5 Local Descent Methods

- Applicable to multivariate problems

- We are searching for a *local model* that provides some guidance in a certain region of f over **where to go to next**.
- Gradient and Hessian are informative about this.

1.5.1 Local Descent Outline

All descent methods follow more or less this structure. At iteration k ,

1. Check if candidate $\mathbf{x}^{(k)}$ satisfies stopping criterion:
 - if yes: stop
 - if no: continue
2. Get the local *descent direction* $\mathbf{d}^{(k)}$, using gradient, hessian, or both.
3. Set the *step size*, i.e. the length of the next step, α^k
4. Get the next candidate via

$$\mathbf{x}^{(k+1)} \leftarrow \mathbf{x}^{(k)} + \alpha^k \mathbf{d}^{(k)}$$

1.5.2 The Line Search Strategy

- An algorithm from the line search class chooses a direction $\mathbf{d}^{(k)} \in \mathbb{R}^n$ and searches along that direction starting from the current iterate $x_k \in \mathbb{R}^n$ for a new iterate $x_{k+1} \in \mathbb{R}^n$ with a lower function value.
- After deciding on a direction $\mathbf{d}^{(k)}$, one needs to decide the *step length* α to travel by solving

$$\min_{\alpha > 0} f(x_k + \alpha \mathbf{d}^{(k)})$$

- In practice, solving this exactly is too costly, so algos usually generate a sequence of trial values α and pick the one with the lowest f .

```
In [4]: # https://github.com/JuliaNLSolvers/LineSearches.jl
using LineSearches
prob = Optim.UnconstrainedProblems.examples["Rosenbrock"]

algo_hz = Newton(linesearch = HagerZhang())
res_hz = Optim.optimize(prob.f, prob.g!, prob.h!, prob.initial_x, method=algo_hz)

UndefVarError: UnconstrainedProblems not defined

Stacktrace:

 [1] include_string(::String, ::String) at ./loading.jl:515
```

1.5.3 The Trust Region Strategy

- First choose max step size, then the direction
- Finds the next step $\mathbf{x}^{(k+1)}$ by minimizing a model of \hat{f} over a *trust region*, centered on $\mathbf{x}^{(k)}$
 - 2nd order Taylor approx of f is common.
- Radius δ of trust region is changed based on how well \hat{f} fits f in trust region.
- Get \mathbf{x}' via

$$\begin{aligned} \min_{\mathbf{x}'} \quad & \hat{f}(\mathbf{x}') \\ \text{subject to} \quad & \|\mathbf{x} - \mathbf{x}'\| \leq \delta \end{aligned}$$

```
In [5]: # Optim.jl has a TrustRegion for Newton (see below for Newton's Method)
        NewtonTrustRegion(; initial_delta = 1.0, # The starting trust region radius
                             delta_hat = 100.0, # The largest allowable trust region radius
                             eta = 0.1, #When rho is at least eta, accept the step.
                             rho_lower = 0.25, # When rho is less than rho_lower, shrink the trust region
                             rho_upper = 0.75) # When rho is greater than rho_upper, grow the trust region
        res = Optim.optimize(prob.f, prob.g!, prob.h!, prob.initial_x, method=NewtonTrustRegion())

        UndefinedVarError: prob not defined

        Stacktrace:

        [1] include_string(::String, ::String) at ./loading.jl:515
```

1.5.4 Stopping criteria

1. maximum number of iterations reached
2. absolute improvement $|f(x) - f(x')| \leq \epsilon$
3. relative improvement $|f(x) - f(x')|/|f(x)| \leq \epsilon$
4. Gradient close to zero $|g(x)| \approx 0$

1.5.5 Gradient Descent

- Here we define

$$\mathbf{g}^{(k)} = \nabla f(\mathbf{d}^{(k)})$$

- And our descent becomes

$$\mathbf{d}^{(k)} = -\nabla \frac{\mathbf{g}^{(k)}}{\|\mathbf{g}^{(k)}\|}$$

- Minimizing wrt step size results in a jagged path (each direction is orthogonal to previous direction!)

$$\alpha^{(k)} = \arg \min \alpha f(\mathbf{x}^{(k)} + \alpha \mathbf{d}^{(k)})$$

- *Conjugate Gradient* avoids this issue.

```

In [6]: # Optim.jl again
        GradientDescent(; alphaguess = LineSearches.InitialPrevious(),
                        linesearch = LineSearches.HagerZhang(),
                        P = nothing,
                        preconditioner = (P, x) -> nothing)

Out[6]: Optim.GradientDescent{LineSearches.InitialPrevious{Float64},LineSearches.HagerZhang{Float64}}
        alpha: Float64 1.0
        alphamin: Float64 0.0
        alphamax: Float64 Inf
        , LineSearches.HagerZhang{Float64}
        delta: Float64 0.1
        sigma: Float64 0.9
        alphamax: Float64 Inf
        rho: Float64 5.0
        epsilon: Float64 1.0e-6
        gamma: Float64 0.66
        linesearchmax: Int64 50
        psi3: Float64 0.1
        display: Int64 0
        , nothing, #3, Optim.Flat()

In [7]: # there is a dedicated LineSearch package: https://github.com/JuliaNLSolvers/LineSearches.jl
        GD = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent())
        GD1 = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent(), Optim.Options{Optim.Options{}})
        GD2 = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent(), Optim.Options{Optim.Options{}})

        println("gradient descent = $GD")
        println("\n")
        println("gradient descent 2 = $GD1")
        println("\n")
        println("gradient descent 3 = $GD2")

        UndefVarError: rosenbrock not defined

        Stacktrace:

        [1] include_string(::String, ::String) at ./loading.jl:515

```

1.6 Second Order Methods

1.6.1 Newton's Method

- We start with a 2nd order Taylor approx over x at step k :

$$q(x) = f(x^{(k)}) + (x - x^{(k)})f'(x^{(k)}) + \frac{(x - x^{(k)})^2}{2}f''(x^{(k)})$$

- We set find it's root and rearrange to find the next step $k + 1$:

$$\frac{\partial q(x)}{\partial x} = f'(x^{(k)}) + (x - x^{(k)})f''(x^{(k)}) = 0$$

$$x^{(k+1)} = x^{(k)} - \frac{f'(x^{(k)})}{f''(x^{(k)})}$$

- The same argument works for multidimensional functions by using Hessian and Gradient
- We would get a descent \mathbf{d}^k by setting:

$$\mathbf{d}^k = -\frac{\mathbf{g}^k}{\mathbf{H}^k}$$

- There are several options to avoid (often costly) computation of the Hessian \mathbf{H} :
 1. Quasi-Newton updates \mathbf{H} starting from identity matrix
 2. Broyden-Fletcher-Goldfarb-Shanno (BFGS) does better with approx linesearch
 3. L-BFGS is the limited memory version for large problems

In [8]: `optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [0.0, 0.0], Newton(), Optim.Options()`

UndefVarError: rosenbrock not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

In [9]: `@show optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [-1.0, 3.0], BFGS());`

UndefVarError: rosenbrock not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

In [10]: `# low memory BFGS`

`@show optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [0.0, 0.0], LBFGS());`

UndefVarError: rosenbrock not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

1.7 Direct Methods

- No derivative information is used - *derivative free*
- If it's very hard / impossible to provide gradient information, this is our only chance.
- Direct methods use other criteria than the gradient to inform the next step (and ultimately convergence).

1.7.1 Cyclic Coordinate Descent -- Taxicab search

- We do a line search over each dimension, one after the other
- *taxicab* because the path looks like a NYC taxi changing direction at each block.
- given $\mathbf{x}^{(1)}$, we proceed

$$\mathbf{x}^{(2)} = \arg \min_{x_1} f(x_1, x_2^{(1)}, \dots, x_n^{(1)})$$

$$\mathbf{x}^{(3)} = \arg \min_{x_2} f(x_1^{(2)}, x_2, x_3^{(2)}, \dots, x_n^{(2)})$$

- unfortunately this can easily get stuck because it can only move in 2 directions.

```
In [11]: # start to setup a basis function, i.e. unit vectors to index each direction:
basis(i, n) = [k == i ? 1.0 : 0.0 for k in 1 : n]
function cyclic_coordinate_descent(f, x, )
    , n = Inf, length(x)
    while abs() >
        x = copy(x)
        for i in 1 : n
            d = basis(i, n)
            x = line_search(f, x, d)
        end
        = norm(x - x)
    end
    return x
end
```

```
Out[11]: cyclic_coordinate_descent (generic function with 1 method)
```

1.7.2 General Pattern Search

- We search according to an arbitrary *pattern* \mathcal{P} of candidate points, anchored at current guess \mathbf{x} .
- With step size α and set \mathcal{D} of directions

$$\mathcal{P} = \mathbf{x} + \alpha \mathbf{d} \text{ for } \mathbf{d} \in \mathcal{D}$$

- Convergence is guaranteed under conditions:

- \mathcal{D} must be a positive spanning set: at least one $\mathbf{d} \in \mathcal{D}$ has a non-zero gradient.

```
In [12]: function generalized_pattern_search(f, x, , D, , =0.5)
    y, n = f(x), length(x)
    evals = 0
    while >
        improved = false
        for (i,d) in enumerate(D)
            x = x + *d
            y = f(x)
            evals += 1
            if y < y
                x, y, improved = x, y, true
                D = unshift!(deleteat!(D, i), d)
                break
            end
        end
        if !improved
            *=
        end
    end
    println("$evals evaluations")
    return x
end

Out[12]: generalized_pattern_search (generic function with 2 methods)

In [13]: D = [[1,0],[0,1],[-1,-0.5]]
    y=generalized_pattern_search(rosenbrock.f,zeros(2),0.8,D,1e-6 )
```

UndefVarError: rosenbrock not defined

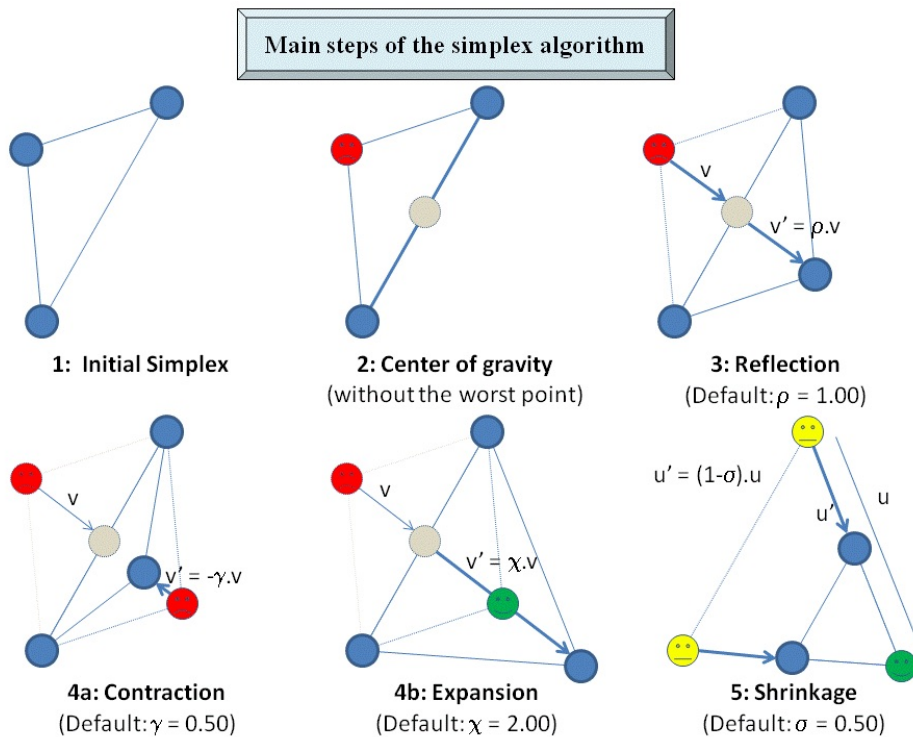
Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

1.8 Bracketing for Multidimensional Problems: Nelder-Mead

- The Goal here is to find the simplex containing the local minimizer x^*
- In the case where f is n -D, this simplex has $n + 1$ vertices
- In the case where f is 2-D, this simplex has $2 + 1$ vertices, i.e. it's a triangle.
- The method proceeds by evaluating the function at all $n + 1$ vertices, and by replacing the worst function value with a new guess.
- this can be achieved by a sequence of moves:

- reflect
- expand
- contract
- shrink movements.



- this is a very popular method. The matlab functions `fmincon` and `fminsearch` implements it.
- When it works, it works quite fast.
- No derivatives required.

In [14]: `optimize(rosenbrock.f, [0.0, 0.0], NelderMead())`

UndefVarError: rosenbrock not defined

Stacktrace:

[1] `include_string(::String, ::String)` at `./loading.jl:515`

- But.

1.9 Bracketing for Multidimensional Problems: Comment on Nelder-Mead

Lagarias et al. (SIOPT, 1999): At present there is no function in any dimension greater than one, for which the original Nelder-Mead algorithm has been proved to converge to a minimizer.

Given all the known inefficiencies and failures of the Nelder-Mead algorithm [...], one might wonder why it is used at all, let alone why it is so extraordinarily popular.

1.10 things to read up on

- Divided Rectangles (DIRECT)
- simulated annealing and other stochastic gradient methods

2 Constraints

2.1 lagrange multipliers

2.2 duality

2.3 penalty methods

2.4 augmented lagrange

2.5 interior point

2.5.1 Examples

$$\min_{x \in \mathbb{R}^2} \sqrt{x_2} \text{ subject to } \begin{aligned} x_2 &\geq 0 \\ x_2 &\geq (a_1 x_1 + b_1)^3 \\ x_2 &\geq (a_2 x_1 + b_2)^3 \end{aligned}$$

2.6 Constrained Optimisation with `NLopt.jl`

- We need to specify one function for each objective and constraint.
- Both of those functions need to compute the function value (i.e. objective or constraint) *and* it's respective gradient.
- `NLopt` expects constraints **always** to be formulated in the format

$$g(x) \leq 0$$

where g is your constraint function

- The constraint function is formulated for each constraint at x . it returns a number (the value of the constraint at x), and it fills out the gradient vector, which is the partial derivative of the current constraint wrt x .
- There is also the option to have vector valued constraints, see the documentation.
- We set this up as follows:

```
In [15]: using NLopt
```

```
count = 0 # keep track of # function evaluations

function myfunc(x::Vector, grad::Vector)
    if length(grad) > 0
        grad[1] = 0
        grad[2] = 0.5/sqrt(x[2])
    end
end
```

```

end

global count
count::Int += 1
println("f_$(count)($x)")

sqrt(x[2])
end

function myconstraint(x::Vector, grad::Vector, a, b)
    if length(grad) > 0
        grad[1] = 3a * (a*x[1] + b)^2
        grad[2] = -1
    end
    (a*x[1] + b)^3 - x[2]
end

opt = Opt(:LD_MMA, 2)
lower_bounds!(opt, [-Inf, 0.])
xtol_rel!(opt, 1e-4)

min_objective!(opt, myfunc)
inequality_constraint!(opt, (x,g) -> myconstraint(x,g,2,0), 1e-8)
inequality_constraint!(opt, (x,g) -> myconstraint(x,g,-1,1), 1e-8)

(minfn,minx,ret) = NLOpt.optimize(opt, [1.234, 5.678])
println("got $minf at $minx after $count iterations (returned $ret)")

```

ArgumentError: Module NLOpt not found in current path.
Run `Pkg.add("NLOpt")` to install the NLOpt package.

Stacktrace:

```

[1] _require(::Symbol) at ./loading.jl:428

[2] require(::Symbol) at ./loading.jl:398

[3] include_string(::String, ::String) at ./loading.jl:515

```

2.7 NLOpt: Rosenbrock

- Let's tackle the rosenbrock example again.
- To make it more interesting, let's add an inequality constraint.

$$\min_{x \in \mathbb{R}^2} (1 - x_1)^2 + 100(x_2 - x_1^2)^2 \text{ subject to } 0.8 - x_1^2 - x_2^2 \geq 0$$

- in NLOpt format, the constraint is $x_1 + x_2 - 0.8 \leq 0$

```
In [16]: function rosenbrockf(x::Vector,grad::Vector)
    if length(grad) > 0
        grad[1] = -2.0 * (1.0 - x[1]) - 400.0 * (x[2] - x[1]^2) * x[1]
        grad[2] = 200.0 * (x[2] - x[1]^2)
    end
    return (1.0 - x[1])^2 + 100.0 * (x[2] - x[1]^2)^2
end
function r_constraint(x::Vector, grad::Vector)
    if length(grad) > 0
        grad[1] = 2*x[1]
        grad[2] = 2*x[2]
    end
    return x[1]^2 + x[2]^2 - 0.8
end
opt = Opt(:LD_MMA, 2)
lower_bounds!(opt, [-5, -5.0])
min_objective!(opt, (x,g) -> rosenbrockf(x,g))
inequality_constraint!(opt, (x,g) -> r_constraint(x,g))
ftol_rel!(opt, 1e-9)
NLOpt.optimize(opt, [-1.0, 0.0])
```

UndefVarError: Opt not defined

Stacktrace:

```
[1] include_string(::String, ::String) at ./loading.jl:515
```

2.8 JuMP.jl

- Introduce [JuMP.jl](#)
- JuMP is a mathematical programming interface for Julia. It is like AMPL, but for free and with a decent programming language.
- The main highlights are:
 - It uses automatic differentiation to compute derivatives from your expression.
 - It supplies this information, as well as the sparsity structure of the Hessian to your preferred solver.
 - It decouples your problem completely from the type of solver you are using. This is great, since you don't have to worry about different solvers having different interfaces.
 - In order to achieve this, JuMP uses [MathProgBase.jl](#), which converts your problem formulation into a standard representation of an optimization problem.
- Let's look at the readme
- The technical citation is Lubin et al [?]

2.9 JuMP: Quick start guide

- this is from the [quick start guide](#)
- please check the docs, they are excellent.

2.9.1 Step 1: create a model

- A model collects variables, objective function and constraints.
- it defines a solver to be used.

```
using Clp
m = Model(solver=ClpSolver()) # provide a solver

# Define variables
@variable(m, x ) # No bounds
@variable(m, x >= lb ) # Lower bound only (note: 'lb <= x' is not valid)
@variable(m, x <= ub ) # Upper bound only
@variable(m, lb <= x <= ub ) # Lower and upper bounds

# we can create arrays of a variable
N = 2
@variable(m, x[1:M,1:N] >= 0 )

# or put them in a block
@variables m begin
    x
    y >= 0
    Z[1:10], Bin
    X[1:3,1:3], SDP
    q[i=1:2], (lowerbound = i, start = 2i, upperbound = 3i)
    t[j=1:3], (Int, start = j)
end

# Equivalent to:
@variable(m, x)
@variable(m, y >= 0)
@variable(m, Z[1:10], Bin)
@variable(m, X[1:3,1:3], SDP)
@variable(m, q[i=1:2], lowerbound = i, start = 2i, upperbound = 3i)
@variable(m, t[j=1:3], Int, start = j)

# bounds can depend on indices
@variable(m, x[i=1:10] >= i )
```

2.10 Objective and Constraints

- We can easily add objective and constraint functions:

```

@constraint(m, x[i] - s[i] <= 0) # Other options: == and >=
@constraint(m, sum(x[i] for i=1:numLocation) == 1)
@objective(m, Max, 5x + 22y + (x+y)/2) # or Min

```

- This is fully integrated with Julia. you can use the generator syntax for sums:

```

@objective(sum(x[i] + y[i]/pi for i = I1, j = I2 if i+j < some_val))

```

```

In [17]: ## Simple example

```

```

using JuMP
using Clp

m = Model(solver = ClpSolver())
@variable(m, 0 <= x <= 2 )
@variable(m, 0 <= y <= 30 )

@objective(m, Max, 5x + 3*y )
@constraint(m, 1x + 5y <= 3.0 )

print(m)

status = solve(m)

println("Objective value: ", getobjectivevalue(m))
println("x = ", getvalue(x))
println("y = ", getvalue(y))

```

```

INFO: Recompiling stale cache file /Users/74097/.julia/lib/v0.6/JuMP.ji for module JuMP.INFO: Pr

```

```

Max 5 x + 3 y
Subject to
  x + 5 y 3
  0 x 2
  0 y 30
Objective value: 10.6
x = 2.0
y = 0.2

```

```

In [18]: # JuMP: Rosenbrock Example

```

```

# Instead of hand-coding first and second derivatives, you only have to give `JuMP` exp
# Here is an example.

```

```

using Ipopt

let

    m = Model(solver=IpoptSolver())

```

```

@variable(m, x)
@variable(m, y)

@NLobjective(m, Min, (1-x)^2 + 100(y-x^2)^2)

solve(m)

println("x = ", getvalue(x), " y = ", getvalue(y))

end

```

INFO: Recompiling stale cache file /Users/74097/.julia/lib/v0.6/Ipopt.ji for module Ipopt.

```

*****
This program contains Ipopt, a library for large-scale nonlinear optimization.
Ipopt is released as open source code under the Eclipse Public License (EPL).
For more information visit http://projects.coin-or.org/Ipopt
*****

```

This is Ipopt version 3.12.9, running with linear solver mumps.
NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

```

Number of nonzeros in equality constraint Jacobian...:      0
Number of nonzeros in inequality constraint Jacobian.:      0
Number of nonzeros in Lagrangian Hessian...:              3

```

```

Total number of variables...:      2
      variables with only lower bounds:      0
      variables with lower and upper bounds:  0
      variables with only upper bounds:      0
Total number of equality constraints...:      0
Total number of inequality constraints...:      0
      inequality constraints with only lower bounds:      0
      inequality constraints with lower and upper bounds:  0
      inequality constraints with only upper bounds:      0

```

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.0000000e+00	0.00e+00	2.00e+00	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	9.5312500e-01	0.00e+00	1.25e+01	-1.0	1.00e+00	-	1.00e+00	2.50e-01f	3
2	4.8320569e-01	0.00e+00	1.01e+00	-1.0	9.03e-02	-	1.00e+00	1.00e+00f	1
3	4.5708829e-01	0.00e+00	9.53e+00	-1.0	4.29e-01	-	1.00e+00	5.00e-01f	2
4	1.8894205e-01	0.00e+00	4.15e-01	-1.0	9.51e-02	-	1.00e+00	1.00e+00f	1
5	1.3918726e-01	0.00e+00	6.51e+00	-1.7	3.49e-01	-	1.00e+00	5.00e-01f	2
6	5.4940990e-02	0.00e+00	4.51e-01	-1.7	9.29e-02	-	1.00e+00	1.00e+00f	1
7	2.9144630e-02	0.00e+00	2.27e+00	-1.7	2.49e-01	-	1.00e+00	5.00e-01f	2
8	9.8586451e-03	0.00e+00	1.15e+00	-1.7	1.10e-01	-	1.00e+00	1.00e+00f	1
9	2.3237475e-03	0.00e+00	1.00e+00	-1.7	1.00e-01	-	1.00e+00	1.00e+00f	1

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
10	2.3797236e-04	0.00e+00	2.19e-01	-1.7	5.09e-02	-	1.00e+00	1.00e+00f	1
11	4.9267371e-06	0.00e+00	5.95e-02	-1.7	2.53e-02	-	1.00e+00	1.00e+00f	1
12	2.8189505e-09	0.00e+00	8.31e-04	-2.5	3.20e-03	-	1.00e+00	1.00e+00f	1
13	1.0095040e-15	0.00e+00	8.68e-07	-5.7	9.78e-05	-	1.00e+00	1.00e+00f	1
14	1.3288608e-28	0.00e+00	2.02e-13	-8.6	4.65e-08	-	1.00e+00	1.00e+00f	1

Number of Iterations....: 14

	(scaled)	(unscaled)
Objective....:	1.3288608467480825e-28	1.3288608467480825e-28
Dual infeasibility....:	2.0183854587685121e-13	2.0183854587685121e-13
Constraint violation....:	0.0000000000000000e+00	0.0000000000000000e+00
Complementarity....:	0.0000000000000000e+00	0.0000000000000000e+00
Overall NLP error....:	2.0183854587685121e-13	2.0183854587685121e-13

Number of objective function evaluations	= 36
Number of objective gradient evaluations	= 15
Number of equality constraint evaluations	= 0
Number of inequality constraint evaluations	= 0
Number of equality constraint Jacobian evaluations	= 0
Number of inequality constraint Jacobian evaluations	= 0
Number of Lagrangian Hessian evaluations	= 14
Total CPU secs in IPOPT (w/o function evaluations)	= 0.117
Total CPU secs in NLP function evaluations	= 0.026

EXIT: Optimal Solution Found.

x = 0.9999999999999989 y = 0.9999999999999792

In [19]: *# not bad, right?*

adding the constraint from before:

let

```
m = Model(solver=IpoptSolver())
```

```
@variable(m, x)
```

```
@variable(m, y)
```

```
@NObjective(m, Min, (1-x)^2 + 100(y-x^2)^2)
```

```
@NLconstraint(m, x^2 + y^2 <= 0.8)
```

```
solve(m)
```

```
println("x = ", getvalue(x), " y = ", getvalue(y))
```

end

This is Ipopt version 3.12.9, running with linear solver mumps.

NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

Number of nonzeros in equality constraint Jacobian...: 0
Number of nonzeros in inequality constraint Jacobian.: 2
Number of nonzeros in Lagrangian Hessian...: 5

Total number of variables...: 2
variables with only lower bounds: 0
variables with lower and upper bounds: 0
variables with only upper bounds: 0

Total number of equality constraints...: 0

Total number of inequality constraints...: 1
inequality constraints with only lower bounds: 0
inequality constraints with lower and upper bounds: 0
inequality constraints with only upper bounds: 1

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.0000000e+00	0.00e+00	2.00e+00	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	9.5312500e-01	0.00e+00	1.25e+01	-1.0	5.00e-01	-	1.00e+00	5.00e-01f	2
2	4.9204994e-01	0.00e+00	9.72e-01	-1.0	8.71e-02	-	1.00e+00	1.00e+00f	1
3	2.0451702e+00	0.00e+00	3.69e+01	-1.7	3.80e-01	-	1.00e+00	1.00e+00H	1
4	1.0409466e-01	0.00e+00	3.10e-01	-1.7	1.46e-01	-	1.00e+00	1.00e+00f	1
5	8.5804626e-02	0.00e+00	2.71e-01	-1.7	9.98e-02	-	1.00e+00	1.00e+00h	1
6	9.4244879e-02	0.00e+00	6.24e-02	-1.7	3.74e-02	-	1.00e+00	1.00e+00h	1
7	8.0582034e-02	0.00e+00	1.51e-01	-2.5	6.41e-02	-	1.00e+00	1.00e+00h	1
8	7.8681242e-02	0.00e+00	2.12e-03	-2.5	1.12e-02	-	1.00e+00	1.00e+00h	1
9	7.6095770e-02	0.00e+00	6.16e-03	-3.8	1.37e-02	-	1.00e+00	1.00e+00h	1
iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
10	7.6033892e-02	0.00e+00	2.23e-06	-3.8	3.99e-04	-	1.00e+00	1.00e+00h	1
11	7.5885642e-02	0.00e+00	2.07e-05	-5.7	7.99e-04	-	1.00e+00	1.00e+00h	1
12	7.5885428e-02	0.00e+00	2.74e-11	-5.7	1.38e-06	-	1.00e+00	1.00e+00h	1
13	7.5883585e-02	0.00e+00	3.19e-09	-8.6	9.93e-06	-	1.00e+00	1.00e+00f	1

Number of Iterations...: 13

	(scaled)	(unscaled)
Objective...:	7.5883585442440671e-02	7.5883585442440671e-02
Dual infeasibility...:	3.1949178858070582e-09	3.1949178858070582e-09
Constraint violation...:	0.0000000000000000e+00	0.0000000000000000e+00
Complementarity...:	2.5454985882932001e-09	2.5454985882932001e-09
Overall NLP error...:	3.1949178858070582e-09	3.1949178858070582e-09

Number of objective function evaluations = 20
Number of objective gradient evaluations = 14

Number of equality constraint evaluations	=	0
Number of inequality constraint evaluations	=	20
Number of equality constraint Jacobian evaluations	=	0
Number of inequality constraint Jacobian evaluations	=	14
Number of Lagrangian Hessian evaluations	=	13
Total CPU secs in IPOPT (w/o function evaluations)	=	0.005
Total CPU secs in NLP function evaluations	=	0.002

EXIT: Optimal Solution Found.

x = 0.7247018392092258 y = 0.5242206029480763

2.11 JuMP: Maximum Likelihood

- Let's redo the maximum likelihood example in JuMP.
- Let μ, σ^2 be the unknown mean and variance of a random sample generated from the normal distribution.
- Find the maximum likelihood estimator for those parameters!
- density:

$$f(x_i|\mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$$

- Likelihood Function

$$\begin{aligned} L(\mu, \sigma^2) &= \prod_{i=1}^N f(x_i|\mu, \sigma^2) = \frac{1}{(\sigma\sqrt{2\pi})^n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2\right) \\ &= (\sigma^2 2\pi)^{-\frac{n}{2}} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2\right) \end{aligned}$$

- Constraints: $\mu \in \mathbb{R}, \sigma > 0$
- log-likelihood:

$$\log L = l = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2$$

- Let's do this in JuMP.

```
In [20]: # Copyright 2015, Iain Dunning, Joey Huchette, Miles Lubin, and contributors
# example modified
using Distributions

distrib = Normal(4.5, 3.5)
n = 10000
```

```

data = rand(distrib,n);

m = Model(solver=IpoptSolver())

@variable(m, mu, start = 0.0)
@variable(m, sigma >= 0.0, start = 1.0)

@NLObjective(m, Max, -(n/2)*log(2*sigma^2)-sum((data[i] - mu) ^ 2 for i = 1:n)/(2*sigma

solve(m)
println(" = ", getvalue(mu),", mean(data) = ", mean(data))
println("^2 = ", getvalue(sigma)^2, ", var(data) = ", var(data))

```

This is Ipopt version 3.12.9, running with linear solver mumps.

NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

```

Number of nonzeros in equality constraint Jacobian...:      0
Number of nonzeros in inequality constraint Jacobian.:      0
Number of nonzeros in Lagrangian Hessian...:           3

```

```

Total number of variables...:      2
      variables with only lower bounds:      1
      variables with lower and upper bounds:    0
      variables with only upper bounds:      0
Total number of equality constraints...:      0
Total number of inequality constraints...:      0
      inequality constraints with only lower bounds:      0
      inequality constraints with lower and upper bounds:    0
      inequality constraints with only upper bounds:      0

```

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.6887292e+05	0.00e+00	1.01e+02	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	1.2430166e+05	0.00e+00	1.03e+02	-1.0	9.54e+00	-	1.00e+00	5.00e-01f	2
2	7.6133884e+04	0.00e+00	4.27e+01	-1.0	2.30e-01	-	8.90e-01	1.00e+00f	1
3	5.0257956e+04	0.00e+00	1.77e+01	-1.0	2.99e-01	-	1.00e+00	1.00e+00f	1
4	3.6881104e+04	0.00e+00	7.15e+00	-1.0	3.78e-01	-	1.00e+00	1.00e+00f	1
5	3.0407076e+04	0.00e+00	2.81e+00	-1.0	4.59e-01	-	1.00e+00	1.00e+00f	1
6	2.7667909e+04	0.00e+00	1.02e+00	-1.0	5.11e-01	-	1.00e+00	1.00e+00f	1
7	2.6785732e+04	0.00e+00	3.17e-01	-1.0	4.84e-01	-	1.00e+00	1.00e+00f	1
8	2.6637016e+04	0.00e+00	5.63e-02	-1.7	3.00e-01	-	1.00e+00	1.00e+00f	1
9	2.6629671e+04	0.00e+00	3.26e-03	-2.5	8.74e-02	-	1.00e+00	1.00e+00f	1
iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
10	2.6629653e+04	0.00e+00	8.70e-06	-3.8	4.73e-03	-	1.00e+00	1.00e+00f	1
11	2.6629653e+04	0.00e+00	1.49e-09	-5.7	6.18e-05	-	1.00e+00	1.00e+00f	1
12	2.6629653e+04	0.00e+00	4.01e-13	-8.6	9.87e-07	-	1.00e+00	1.00e+00f	1

Number of Iterations...: 12

	(scaled)	(unscaled)
Objective....:	8.6077852994465989e+00	2.6629653293957461e+04
Dual infeasibility...:	4.0102196914099508e-13	1.2406299216328470e-09
Constraint violation...:	0.0000000000000000e+00	0.0000000000000000e+00
Complementarity...:	2.5064286232902822e-09	7.7540648287966047e-06
Overall NLP error...:	2.5064286232902822e-09	7.7540648287966047e-06

Number of objective function evaluations	=	18
Number of objective gradient evaluations	=	13
Number of equality constraint evaluations	=	0
Number of inequality constraint evaluations	=	0
Number of equality constraint Jacobian evaluations	=	0
Number of inequality constraint Jacobian evaluations	=	0
Number of Lagrangian Hessian evaluations	=	12
Total CPU secs in IPOPT (w/o function evaluations)	=	0.005
Total CPU secs in NLP function evaluations	=	0.017

EXIT: Optimal Solution Found.

= 4.460816556166759, mean(data) = 4.46081655616676
 $\hat{2}$ = 12.037822802323896, var(data) = 12.039026695664448

In []: