

SciencesPo Computational Economics Spring 2017

Florian Oswald

January 17, 2019

1 Optimization 2: Algorithms and Constraints

Florian Oswald Sciences Po, 2018

1.1 Bracketing

- A derivative-free method for *univariate* f
- works only on **unimodal** f
- (Draw choosing initial points and where to move next)

1.2 The Golden Ratio or Bracketing Search for 1D problems

- A derivative-free method
- a Bracketing method
 - find the local minimum of f on $[a, b]$
 - select 2 interior points c, d such that $a < c < d < b$
 - * $f(c) \leq f(d) \implies$ min must lie in $[a, d]$. replace b with d , start again with $[a, d]$
 - * $f(c) > f(d) \implies$ min must lie in $[c, b]$. replace a with c , start again with $[c, b]$
 - how to choose b, d though?
 - we want the length of the interval to be independent of whether we replace upper or lower bound
 - we want to reuse the non-replaced point from the previous iteration.
 - these imply the golden rule:
 - new point $x_i = a + \alpha_i(b - a)$, where $\alpha_1 = \frac{3-\sqrt{5}}{2}, \alpha_2 = \frac{\sqrt{5}-1}{2}$
 - α_2 is known as the *golden ratio*, well known for it's role in renaissance art.

```
In [1]: using Plots
        using Optim
        plotlyjs()
        f(x) = exp(x) - x^4
        minf(x) = -f(x)
        brent = optimize(minf, 0, 2, Brent())
        golden = optimize(minf, 0, 2, GoldenSection())
```

```

println("brent = $brent")
println("golden = $golden")
plot(f,0,2)

brent = Results of Optimization Algorithm
* Algorithm: Brent's Method
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 12
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 13
golden = Results of Optimization Algorithm
* Algorithm: Golden Section Search
* Search Interval: [0.000000, 2.000000]
* Minimizer: 8.310315e-01
* Minimum: -1.818739e+00
* Iterations: 37
* Convergence: max(|x - x_upper|, |x - x_lower|) <= 2*(1.5e-08*|x|+2.2e-16): true
* Objective Function Calls: 38

```

1.2.1 Bisection Methods

- Root finding: `Roots.jl`
- Root finding in multivariate functions: `IntervalRootFinding.jl`

```

In [2]: using Roots
        #~find the zeros of this function:
        f(x) = exp(x) - x^4
        ## bracketing
        fzero(f, 8, 9)      # 8.613169456441398
        fzero(f, -10, 0) # -0.8155534188089606

```

```
Out[2]: -0.8155534188089606
```

```

In [3]: using IntervalRootFinding, IntervalArithmetic
        -10..10

```

ArgumentError: Module IntervalRootFinding not found in current path.
 Run `Pkg.add("IntervalRootFinding")` to install the IntervalRootFinding package.

Stacktrace:

```
[1] _require(::Symbol) at ./loading.jl:435
```

```
[2] require(::Symbol) at ./loading.jl:405
```

```
[3] include_string(::String, ::String) at ./loading.jl:522
```

```
In [4]: X = IntervalBox(1..3, 2..4)
```

```
UndefVarError: IntervalBox not defined
```

```
Stacktrace:
```

```
[1] include_string(::String, ::String) at ./loading.jl:522
```

```
In [5]: a = @interval(0.1, 0.3)
       b = @interval(0.3, 0.6)
       a + b
```

```
UndefVarError: @interval not defined
```

```
Stacktrace:
```

```
[1] include_string(::String, ::String) at ./loading.jl:522
```

```
In [6]: rts = IntervalRootFinding.roots(x->x^2 - 2, -10..10, Bisection)
```

```
UndefVarError: IntervalRootFinding not defined
```

```
Stacktrace:
```

```
[1] include_string(::String, ::String) at ./loading.jl:522
```

1.3 Rosenbrock Banana and Optim.jl

- We can supply the objective function and - depending on the solution algorithm - the gradient and hessian as well.

```
In [7]: using Optim
        using OptimTestProblems
        for (name, prob) in MultivariateProblems.UnconstrainedProblems.examples
            println(name)
        end
```

```
Rosenbrock
Quadratic Diagonal
Hosaki
Large Polynomial
Penalty Function I
Beale
Extended Rosenbrock
Polynomial
Powell
Exponential
Paraboloid Diagonal
Paraboloid Random Matrix
Extended Powell
Trigonometric
Fletcher-Powell
Parabola
Himmelblau
```

```
In [8]: rosenbrock = MultivariateProblems.UnconstrainedProblems.examples["Rosenbrock"]
```

```
Out[8]: OptimTestProblems.MultivariateProblems.OptimizationProblem{Void,Void,Float64,String,Void}
```

1.4 Comparison Methods

- We will now look at a first class of algorithms, which are very simple, but sometimes a good starting point.
- They just *compare* function values.
- *Grid Search* : Compute the objective function at $G = \{x_1, \dots, x_N\}$ and pick the highest value of f .
 - This is very slow.
 - It requires large N .
 - But it's robust (will find global optimizer for large enough N)

```
In [9]: # grid search on rosenbrock
        grid = collect(-1.0:0.1:3);
        grid2D = [[i;j] for i in grid,j in grid];
        val2D = map(rosenbrock.f,grid2D);
        r = findmin(val2D);
        println("grid search results in minimizer = $(grid2D[r[2]])")
```

```
grid search results in minimizer = [1.0, 1.0]
```

1.5 Local Descent Methods

- Applicable to multivariate problems
- We are searching for a *local model* that provides some guidance in a certain region of f over **where to go to next**.
- Gradient and Hessian are informative about this.

1.5.1 Local Descent Outline

All descent methods follow more or less this structure. At iteration k ,

1. Check if candidate $\mathbf{x}^{(k)}$ satisfies stopping criterion:
 - if yes: stop
 - if no: continue
2. Get the local *descent direction* $\mathbf{d}^{(k)}$, using gradient, hessian, or both.
3. Set the *step size*, i.e. the length of the next step, α^k
4. Get the next candidate via

$$\mathbf{x}^{(k+1)} \leftarrow \mathbf{x}^{(k)} + \alpha^k \mathbf{d}^{(k)}$$

1.5.2 The Line Search Strategy

- An algorithm from the line search class chooses a direction $\mathbf{d}^{(k)} \in \mathbb{R}^n$ and searches along that direction starting from the current iterate $x_k \in \mathbb{R}^n$ for a new iterate $x_{k+1} \in \mathbb{R}^n$ with a lower function value.
- After deciding on a direction $\mathbf{d}^{(k)}$, one needs to decide the *step length* α to travel by solving

$$\min_{\alpha > 0} f(x_k + \alpha \mathbf{d}^{(k)})$$

- In practice, solving this exactly is too costly, so algos usually generate a sequence of trial values α and pick the one with the lowest f .

```
In [10]: # https://github.com/JuliaNLSolvers/LineSearches.jl
using LineSearches
```

```
algo_hz = Newton(linesearch = HagerZhang())
res_hz = Optim.optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, rosenbrock.initial_x, algo_hz)
```

Out[10]: Results of Optimization Algorithm

```
* Algorithm: Newton's Method
* Starting Point: [-1.2,1.0]
* Minimizer: [1.0000000000000033,1.0000000000000067]
* Minimum: 1.109336e-29
* Iterations: 23
* Convergence: true
* |x - x'| 1.0e-32: false
  |x - x'| = 1.13e-08
* |f(x) - f(x')| 1.0e-32 |f(x)|: false
  |f(x) - f(x')| = 6.35e+13 |f(x)|
```

```

* |g(x)| 1.0e-08: true
  |g(x)| = 6.66e-15
* Stopped by an increasing objective: false
* Reached Maximum Number of Iterations: false
* Objective Calls: 71
* Gradient Calls: 71
* Hessian Calls: 23

```

1.5.3 The Trust Region Strategy

- First choose max step size, then the direction
- Finds the next step $\mathbf{x}^{(k+1)}$ by minimizing a model of \hat{f} over a *trust region*, centered on $\mathbf{x}^{(k)}$
 - 2nd order Taylor approx of f is common.
- Radius δ of trust region is changed based on how well \hat{f} fits f in trust region.
- Get \mathbf{x}' via

$$\begin{aligned} \min_{\mathbf{x}'} \quad & \hat{f}(\mathbf{x}') \\ \text{subject to} \quad & \|\mathbf{x} - \mathbf{x}'\| \leq \delta \end{aligned}$$

```

In [11]: # Optim.jl has a TrustRegion for Newton (see below for Newton's Method)
NewtonTrustRegion(; initial_delta = 1.0, # The starting trust region radius
  delta_hat = 100.0, # The largest allowable trust region radius
  eta = 0.1, #When rho is at least eta, accept the step.
  rho_lower = 0.25, # When rho is less than rho_lower, shrink the t
  rho_upper = 0.75) # When rho is greater than rho_upper, grow the
res = Optim.optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, rosenbrock.initial_x

```

Out[11]: Results of Optimization Algorithm

```

* Algorithm: Newton's Method (Trust Region)
* Starting Point: [-1.2,1.0]
* Minimizer: [0.9999999994405535,0.99999999988644926]
* Minimum: 3.405841e-19
* Iterations: 25
* Convergence: true
* |x - x'| 1.0e-32: false
  |x - x'| = 8.84e-06
* |f(x) - f(x')| 1.0e-32 |f(x)|: false
  |f(x) - f(x')| = 1.87e+08 |f(x)|
* |g(x)| 1.0e-08: true
  |g(x)| = 5.53e-09
* Stopped by an increasing objective: false
* Reached Maximum Number of Iterations: false
* Objective Calls: 26
* Gradient Calls: 26
* Hessian Calls: 22

```

1.5.4 Stopping criteria

1. maximum number of iterations reached

2. absolute improvement $|f(x) - f(x')| \leq \epsilon$
3. relative improvement $|f(x) - f(x')|/|f(x)| \leq \epsilon$
4. Gradient close to zero $|g(x)| \approx 0$

1.5.5 Gradient Descent

- Here we define

$$\mathbf{g}^{(k)} = \nabla f(\mathbf{d}^{(k)})$$

- And our descent becomes

$$\mathbf{d}^{(k)} = -\nabla \frac{\mathbf{g}^{(k)}}{\|\mathbf{g}^{(k)}\|}$$

- Minimizing wrt step size results in a jagged path (each direction is orthogonal to previous direction!)

$$\alpha^{(k)} = \arg \min \alpha f(\mathbf{x}^{(k)} + \alpha \mathbf{d}^{(k)})$$

- *Conjugate Gradient* avoids this issue.

In [12]: *# Optim.jl again*

```
GradientDescent(; alphaguess = LineSearches.InitialPrevious(),
                 linesearch = LineSearches.HagerZhang(),
                 P = nothing,
                 preconditioner = (P, x) -> nothing)
```

Out [12]: Optim.GradientDescent{LineSearches.InitialPrevious{Float64},LineSearches.HagerZhang{F

```
  alpha: Float64 1.0
  alphamin: Float64 0.0
  alphamax: Float64 Inf
, LineSearches.HagerZhang{Float64}
  delta: Float64 0.1
  sigma: Float64 0.9
  alphamax: Float64 Inf
  rho: Float64 5.0
  epsilon: Float64 1.0e-6
  gamma: Float64 0.66
  linesearchmax: Int64 50
  psi3: Float64 0.1
  display: Int64 0
, nothing, #5, Optim.Flat())
```

In [13]: *# there is a dedicated LineSearch package: <https://github.com/JuliaNLSolvers/LineSearch>*

```
GD = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent())
GD1 = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent(), Optim.Options{
GD2 = optimize(rosenbrock.f, rosenbrock.g!, [0.0, 0.0], GradientDescent(), Optim.Options{

println("gradient descent = $GD")
println("\n")
println("gradient descent 2 = $GD1")
println("\n")
println("gradient descent 3 = $GD2")
```

gradient descent = Results of Optimization Algorithm

- * Algorithm: Gradient Descent
- * Starting Point: [0.0,0.0]
- * Minimizer: [0.9356732500354086,0.875073922357589]
- * Minimum: 4.154782e-03
- * Iterations: 1000
- * Convergence: false
 - * $|x - x'|$ 1.0e-32: false
 - $|x - x'| = 1.82e-04$
 - * $|f(x) - f(x')|$ 1.0e-32 $|f(x)|$: false
 - $|f(x) - f(x')| = 1.97e-03 |f(x)|$
 - * $|g(x)|$ 1.0e-08: false
 - $|g(x)| = 8.21e-02$
 - * Stopped by an increasing objective: false
 - * Reached Maximum Number of Iterations: true
- * Objective Calls: 2532
- * Gradient Calls: 2532

gradient descent 2 = Results of Optimization Algorithm

- * Algorithm: Gradient Descent
- * Starting Point: [0.0,0.0]
- * Minimizer: [0.9978398797724763,0.9956717950747302]
- * Minimum: 4.682073e-06
- * Iterations: 5000
- * Convergence: false
 - * $|x - x'|$ 1.0e-32: false
 - $|x - x'| = 5.08e-06$
 - * $|f(x) - f(x')|$ 1.0e-32 $|f(x)|$: false
 - $|f(x) - f(x')| = 1.62e-03 |f(x)|$
 - * $|g(x)|$ 1.0e-08: false
 - $|g(x)| = 2.53e-03$
 - * Stopped by an increasing objective: false
 - * Reached Maximum Number of Iterations: true
- * Objective Calls: 12532
- * Gradient Calls: 12532

gradient descent 3 = Results of Optimization Algorithm

- * Algorithm: Gradient Descent
- * Starting Point: [0.0,0.0]
- * Minimizer: [0.9999999914304203,0.9999999828109042]
- * Minimum: 7.368706e-17
- * Iterations: 20458
- * Convergence: true
 - * $|x - x'|$ 1.0e-32: false
 - $|x - x'| = 2.00e-11$
 - * $|f(x) - f(x')|$ 1.0e-32 $|f(x)|$: false


```

|f(x) - f(x')| = 1.61e-03 |f(x)|
* |g(x)| 1.0e-08: true
  |g(x)| = 9.99e-09
* Stopped by an increasing objective: false
* Reached Maximum Number of Iterations: false
* Objective Calls: 51177
* Gradient Calls: 51177

```

1.6 Second Order Methods

1.6.1 Newton's Method

- We start with a 2nd order Taylor approx over x at step k :

$$q(x) = f(x^{(k)}) + (x - x^{(k)})f'(x^{(k)}) + \frac{(x - x^{(k)})^2}{2}f''(x^{(k)})$$

- We set find it's root and rearrange to find the next step $k + 1$:

$$\frac{\partial q(x)}{\partial x} = f'(x^{(k)}) + (x - x^{(k)})f''(x^{(k)}) = 0$$

$$x^{(k+1)} = x^{(k)} - \frac{f'(x^{(k)})}{f''(x^{(k)})}$$

- The same argument works for multidimensional functions by using Hessian and Gradient
- We would get a descent \mathbf{d}^k by setting:

$$\mathbf{d}^k = -\frac{\mathbf{g}^k}{\mathbf{H}^k}$$

- There are several options to avoid (often costly) computation of the Hessian \mathbf{H} :
 1. Quasi-Newton updates \mathbf{H} starting from identity matrix
 2. Broyden-Fletcher-Goldfarb-Shanno (BFGS) does better with approx linesearch
 3. L-BFGS is the limited memory version for large problems

In [14]: `optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [0.0, 0.0], Newton(), Optim.Options{`

Iter	Function value	Gradient norm
0	1.000000e+00	2.000000e+00
1	8.431140e-01	1.588830e+00
2	6.776980e-01	3.453340e+00
3	4.954645e-01	4.862093e+00
4	3.041921e-01	2.590086e+00
5	1.991512e-01	3.780900e+00
6	9.531907e-02	1.299090e+00
7	5.657827e-02	2.445401e+00
8	2.257807e-02	1.839332e+00
9	6.626125e-03	1.314236e+00
10	8.689753e-04	5.438279e-01

11	4.951399e-06	7.814556e-02
12	9.065070e-10	6.017046e-04
13	9.337686e-18	1.059738e-07
14	3.081488e-31	1.110223e-15

Out[14]: Results of Optimization Algorithm

```

* Algorithm: Newton's Method
* Starting Point: [0.0,0.0]
* Minimizer: [0.9999999999999994,0.9999999999999989]
* Minimum: 3.081488e-31
* Iterations: 14
* Convergence: true
  * |x - x'| 1.0e-32: false
    |x - x'| = 3.06e-09
  * |f(x) - f(x')| 1.0e-32 |f(x)|: false
    |f(x) - f(x')| = 3.03e+13 |f(x)|
  * |g(x)| 1.0e-08: true
    |g(x)| = 1.11e-15
  * Stopped by an increasing objective: false
  * Reached Maximum Number of Iterations: false
* Objective Calls: 44
* Gradient Calls: 44
* Hessian Calls: 14

```

In [15]: @show optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [-1.0, 3.0], BFGS());

optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [-1.0, 3.0], BFGS()) = Results of Optimiz

```

* Algorithm: BFGS
* Starting Point: [-1.0,3.0]
* Minimizer: [0.9999999999999956,0.999999999999987]
* Minimum: 1.707144e-27
* Iterations: 39
* Convergence: true
  * |x - x'| 1.0e-32: false
    |x - x'| = 1.54e-08
  * |f(x) - f(x')| 1.0e-32 |f(x)|: false
    |f(x) - f(x')| = 3.55e+10 |f(x)|
  * |g(x)| 1.0e-08: true
    |g(x)| = 1.63e-12
  * Stopped by an increasing objective: false
  * Reached Maximum Number of Iterations: false
* Objective Calls: 137
* Gradient Calls: 137

```

In [16]: # low memory BFGS

@show optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [0.0, 0.0], LBFGS());

```

optimize(rosenbrock.f, rosenbrock.g!, rosenbrock.h!, [0.0, 0.0], LBFGS()) = Results of Optimiz
* Algorithm: L-BFGS
* Starting Point: [0.0,0.0]
* Minimizer: [0.999999999999928,0.9999999999998559]
* Minimum: 5.191703e-27
* Iterations: 24
* Convergence: true
* |x - x'| 1.0e-32: false
  |x - x'| = 4.58e-11
* |f(x) - f(x')| 1.0e-32 |f(x)|: false
  |f(x) - f(x')| = 8.50e+07 |f(x)|
* |g(x)| 1.0e-08: true
  |g(x)| = 1.44e-13
* Stopped by an increasing objective: false
* Reached Maximum Number of Iterations: false
* Objective Calls: 67
* Gradient Calls: 67

```

Direct Methods

- No derivative information is used - *derivative free*
- If it's very hard / impossible to provide gradient information, this is our only chance.
- Direct methods use other criteria than the gradient to inform the next step (and ultimately convergence).

1.6.2 Cyclic Coordinate Descent – Taxicab search

- We do a line search over each dimension, one after the other
- *taxicab* because the path looks like a NYC taxi changing direction at each block.
- given $\mathbf{x}^{(1)}$, we proceed

$$\mathbf{x}^{(2)} = \arg \min_{x_1} f(x_1, x_2^{(1)}, \dots, x_n^{(1)})$$

$$\mathbf{x}^{(3)} = \arg \min_{x_2} f(x_1^{(2)}, x_2, x_3^{(2)}, \dots, x_n^{(2)})$$

- unfortunately this can easily get stuck because it can only move in 2 directions.

```

In [17]: # start to setup a basis function, i.e. unit vectors to index each direction:
basis(i, n) = [k == i ? 1.0 : 0.0 for k in 1 : n]
function cyclic_coordinate_descent(f, x, )
    , n = Inf, length(x)
    while abs() >
        x = copy(x)
        for i in 1 : n
            d = basis(i, n)
            x = line_search(f, x, d)
        end
        = norm(x - x)
    end
end

```

```

        return x
    end

```

Out[17]: cyclic_coordinate_descent (generic function with 1 method)

1.6.3 General Pattern Search

- We search according to an arbitrary *pattern* \mathcal{P} of candidate points, anchored at current guess \mathbf{x} .
- With step size α and set \mathcal{D} of directions

$$\mathcal{P} = \mathbf{x} + \alpha \mathbf{d} \text{ for } \mathbf{d} \in \mathcal{D}$$

- Convergence is guaranteed under conditions:
 - \mathcal{D} must be a positive spanning set: at least one $\mathbf{d} \in \mathcal{D}$ has a non-zero gradient.

```

In [18]: function generalized_pattern_search(f, x, , D, , =0.5)
    y, n = f(x), length(x)
    evals = 0
    while >
        improved = false
        for (i,d) in enumerate(D)
            x = x + *d
            y = f(x)
            evals += 1
            if y < y
                x, y, improved = x, y, true
                D = unshift!(deleteat!(D, i), d)
                break
            end
        end
        if !improved
            *=
        end
    end
    println("$evals evaluations")
    return x
end

```

Out[18]: generalized_pattern_search (generic function with 2 methods)

```

In [19]: D = [[1,0],[0,1],[-1,-0.5]]
    D = [[1,0],[0,1]]
    y=generalized_pattern_search(rosenbrock.f,zeros(2),0.8,D,1e-6 )

```

11923 evaluations

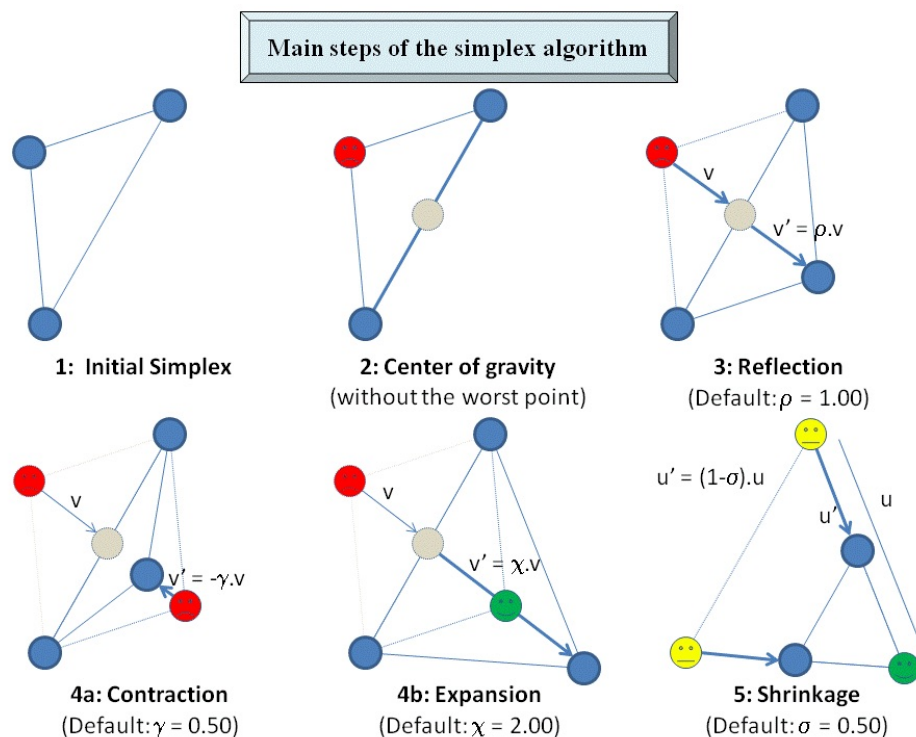
```

Out[19]: 2-element Array{Float64,1}:
 0.999673
 0.999347

```

1.7 Bracketing for Multidimensional Problems: Nelder-Mead

- The Goal here is to find the simplex containing the local minimizer x^*
- In the case where f is n -D, this simplex has $n + 1$ vertices
- In the case where f is 2-D, this simplex has $2 + 1$ vertices, i.e. it's a triangle.
- The method proceeds by evaluating the function at all $n + 1$ vertices, and by replacing the worst function value with a new guess.
- this can be achieved by a sequence of moves:
 - reflect
 - expand
 - contract
 - shrink movements.



- this is a very popular method. The matlab functions `fmincon` and `fminsearch` implements it.
- When it works, it works quite fast.
- No derivatives required.

In [20]: `nm=optimize(rosenbrock.f, [0.0, 0.0], NelderMead());`
`nm.minimizer`

Out [20]: 2-element Array{Float64,1}:
 0.999963
 0.999932

- But.

1.8 Bracketing for Multidimensional Problems: Comment on Nelder-Mead

Lagarias et al. (SIOPT, 1999): At present there is no function in any dimension greater than one, for which the original Nelder-Mead algorithm has been proved to converge to a minimizer.

Given all the known inefficiencies and failures of the Nelder-Mead algorithm [...], one might wonder why it is used at all, let alone why it is so extraordinarily popular.

1.9 things to read up on

- Divided Rectangles (DIRECT)
- simulated annealing and other stochastic gradient methods

1.10 Stochastic Optimization Methods

- Gradient based methods like steepest descent may be susceptible to getting stuck at local minima.
- Randomly shocking the value of the descent direction may be a solution to this.
- For example, one could modify our gradient descent from before to become

$$\mathbf{x}^{(k+1)} \leftarrow \mathbf{x}^{(k)} + \alpha^k \mathbf{g}^{(k)} + \mathbf{u}^{(k)}$$

- where $\mathbf{u}^{(k)} \sim N(0, \sigma_k^2)$, decreasing with k .
- This *stochastic gradient descent* is often used when training neural networks.

1.10.1 Simulated Annealing

- We specify a *temperature* that controls the degree of randomness.
- At first the temperature is high, letting the search jump around widely. This is to escape local minima.
- The temperature is gradually decreased, reducing the step sizes. This is to find the local optimum in the *best* region.
- At every iteration k , we accept new point \mathbf{x}' with

$$\Pr(\text{accept } \mathbf{x}') = \begin{cases} 1 & \text{if } \Delta y \leq 0 \\ \min(e^{\Delta y/t}, 1) & \text{if } \Delta y > 0 \end{cases}$$

- here $\Delta y = f(\mathbf{x}') - f(\mathbf{x})$, and t is the *temperature*.
- $\Pr(\text{accept } \mathbf{x}')$ is called the **Metropolis Criterion**, building block of *Accept/Reject* algorithms.

```
In [21]: #ãf: function
        # x: initial point
        # T: transition distribution
        #ãt: temp schedule, k_max: max iterations
        function simulated_annealing(f, x, T, t, k_max)
            y = f(x)
            ytrace = zeros(typeof(y),k_max)
```

```

x_best, y_best = x, y
for k in 1 : k_max
    x = x + rand(T)
    y = f(x)
    y = y - y
    if y > 0 || rand() < exp(-y/t(k))
        x, y = x, y
    end
    if y < y_best
        x_best, y_best = x, y
    end
    ytrace[k] = y_best
end
return x_best, ytrace
end

```

Out[21]: simulated_annealing (generic function with 1 method)

```

In [22]: function ackley(x, a=20, b=0.2, c=2)
    d = length(x)
    return -a*exp(-b*sqrt(sum(x.^2)/d)) - exp(sum(cos.(c*xi) for xi in x)/d) + a + e
end
using Plots
plotlyjs()
surface(-30:0.1:30,-30:0.1:30,(x,y)->ackley([x, y]))

```

```

In [23]: using Distributions
d = Dict()
for sig in (1,5,25), t1 in (1,10,25)
    tmp = [simulated_annealing(ackley,[15,15],MvNormal(2,sig),x->t1/x,100) for i in 1:d]
    d[(sig,t1)] = Dict()
    d[(sig,t1)][:y] = mapslices(x->ackley(x),hcat([tmp[i][1] for i in 1:300]...),[1])
    d[(sig,t1)][:ytrace] = hcat([tmp[i][2] for i in 1:300]...)
end
d

```

```

# x=[simulated_annealing(ackley,[15,15],MuNormal(2,1),x->1.0/x,100) for i in 1:100]
# y=[simulated_annealing(ackley,[15,15],MuNormal(2,5),x->10.0/x,100) for i in 1:100]
# map((x)->ackley([x[1],x[2]]),y)

```

Out[23]: Dict{Any,Any} with 9 entries:

```

(5, 25) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(5, 10) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(1, 25) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(5, 1)  => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(25, 1) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(25, 25) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(1, 1)  => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(1, 10) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1
(25, 10) => Dict{Any,Any}(Pair{Any,Any}(:ytrace, [19.0043 19.0043 19.0043 1

```

2 Constraints

Recall our core optimization problem:

$$\min_{x \in \mathbb{R}^n} f(x) \text{ s.t. } x \in \mathcal{X}$$

- Up to now, the feasible set was $\mathcal{X} \in \mathbb{R}^n$.
- In **constrained problems** \mathcal{X} is a subset thereof.
- We already encountered *box constraints*, e.g. $x \in [a, b]$.
- Sometimes the constrained solution coincides with the unconstrained one, sometimes it does not.
- There are *equality constraints* and *inequality constraints*.

2.1 Lagrange Multipliers

- Used to optimize a function subject to equality constraints.

$$\begin{aligned} \min_x f(x) \\ \text{subject to } h(x) = 0 \end{aligned}$$

where both f and h have continuous partial derivatives.

- We look for contour lines of f that are aligned to contours of $h(x) = 0$.

In other words, we want to find the best x s.t. $h(x) = 0$ and we have

$$\nabla f(x) = \lambda \nabla h(x)$$

for some *Lagrange Multiplier* λ * Notice that we need the scalar λ because the magnitudes of the gradients may be different. * We therefore form the the **Lagrangian**:

$$\mathcal{L}(x, \lambda) = f(x) - \lambda h(x)$$

2.1.1 Example

Suppose we have

$$\begin{aligned} \min_x -\exp \left(- \left(x_1 x_2 - \frac{3}{2} \right)^2 - \left(x_2 - \frac{3}{2} \right)^2 \right) \\ \text{subject to } x_1 - x_2^2 = 0 \end{aligned}$$

We form the Lagrangian:

$$\mathcal{L}(x_1, x_2, \lambda) = -\exp \left(- \left(x_1 x_2 - \frac{3}{2} \right)^2 - \left(x_2 - \frac{3}{2} \right)^2 \right) - \lambda (x_1 - x_2^2)$$

Then we compute the gradient wrt to x_1, x_2, λ , set to zero and solve.


```
In [35]: gr()
         f(x1,x2) = -exp.(-(x1.*x2 - 3/2).^2 - (x2-3/2).^2)
         c(x1) = sqrt(x1)
         x=0:0.01:3.5
         contour(x,x,(x,y)->f(x,y),lw=1.5,levels=[collect(0:-0.1:-0.85)...,-0.887,-0.95,-1])
         plot!(c,0.01,3.5,label="",lw=2,color=:black)
         scatter!([1.358],[1.165],markersize=5,markercolor=:red,label="Constr. Optimum")
```

- If we had multiple constraints (l), we'd just add them up to get

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\nu}) = f(\mathbf{x}) - \sum_{i=1}^l \lambda_i h_i(\mathbf{x})$$

2.2 Inequality Constraints

Suppose now we had

$$\begin{aligned} & \min_{\mathbf{x}} f(\mathbf{x}) \\ & \text{subject to } g(\mathbf{x}) \leq 0 \end{aligned}$$

which, if the solution lies right on the constraint *boundary*, means that

$$\nabla f - \mu \nabla g = 0$$

for some scalar μ - as before.

- In this case, we say the **constraint is active**.
- In the opposite case, i.e. the solution lies **inside** the constrained region, we say the constraint is **inactive**.
- In that case, we are back to an *unconstrained* problem, look for $\nabla f = 0$, and set $\mu = 0$.

```
In [38]: #~the blue area shows the FEASIBLE SET
         contour(x,x,(x,y)->f(x,y),lw=1.5,levels=[collect(0:-0.1:-0.85)...,-0.887,-0.95,-1])
         plot!(c,0.01,3.5,label="",lw=2,color=:black,fill=(0,0.5,:blue))
         scatter!([1.358],[1.165],markersize=5,markercolor=:red,label="Constr. Optimum")
```

```
In [39]: #~the blue area shows the FEASIBLE SET
         #~NOW THE CONSTRAINT IS INACTIVE OR SLACK!
         c2(x1) = 1+sqrt(x1)
         contour(x,x,(x,y)->f(x,y),lw=1.5,levels=[collect(0:-0.1:-0.85)...,-0.887,-0.95,-1])
         plot!(c2,0.01,3.5,label="",lw=2,color=:black,fill=(0,0.5,:blue))
         scatter!([1],[1.5],markersize=5,markercolor=:red,label="Unconstr. Optimum")
```

2.3 Infinity Step

- We could do an **infinite step** to avoid *infeasible points*:

$$f_{\infty\text{-step}} = \begin{cases} f(\mathbf{x}) & \text{if } g(\mathbf{x}) \leq 0 \\ \infty & \text{else.} \end{cases}$$

$$= f(\mathbf{x}) + \infty(g(\mathbf{x}) > 0)$$

- Unfortunately, this is discontinuous and non-differentiable, i.e. hard to handle for algorithms.
- Instead, we use a *linear penalty* $\mu g(\mathbf{x})$ on the objective if the constraint is violated.
- The penalty provides a lower bound to ∞ :

$$\mathcal{L}(\mathbf{x}, \mu) = f(\mathbf{x}) + \mu g(\mathbf{x})$$

- We can get back the infinite step by maximizing the penalty:

$$f_{\infty\text{-step}} = \max_{\mu \geq 0} \mathcal{L}(\mathbf{x}, \mu)$$

- Every infeasible \mathbf{x} returns ∞ , all others return $f(\mathbf{x})$

2.4 Kuhn-Karush-Tucker (KKT)

- Our problem thus becomes

$$\min_{\mathbf{x}} \max_{\mu \geq 0} \mathcal{L}(\mathbf{x}, \mu)$$

- This is called the **primal problem**. Optimizing this requires:
 1. $g(\mathbf{x}^*) \leq 0$. Point is feasible.
 2. $\mu \geq 0$. Penalty goes into the right direction. *Dual feasibility*.
 3. $\mu g(\mathbf{x}^*) = 0$. Feasible point on the boundary has $g(\mathbf{x}) = 0$, otherwise $g(\mathbf{x}) < 0$ and $\mu = 0$.
 4. $\nabla f(\mathbf{x}^*) - \mu \nabla g(\mathbf{x}^*) = 0$. With an active constraint, we want parallel contours of objective and constraint. When inactive, our optimum just has $\nabla f(\mathbf{x}^*) = 0$, which means $\mu = 0$.

The preceding four conditions are called the Kuhn-Karush-Tucker (KKT) conditions. In the above order, and in general terms, they are:

1. Feasibility
2. Dual Feasibility
3. Complementary Slackness
4. Stationarity.

The KKT conditions are the FONCs for problems with smooth constraints.

2.5 Duality

We can combine equality and inequality constraints:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) = f(\mathbf{x}) + \sum_i \lambda_i h_i(\mathbf{x}) + \sum_j \mu_j g_j(\mathbf{x})$$

where, notice, we reverted the sign of λ since this is unrestricted.

- The Primal problem is identical to the original problem and just as difficult to solve:

$$\min_{\mathbf{x}} \max_{\boldsymbol{\lambda} \geq 0, \boldsymbol{\mu}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu})$$

- The Dual problem reverses min and max:

$$\max_{\boldsymbol{\lambda} \geq 0, \boldsymbol{\mu}} \min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu})$$

2.5.1 Dual Values

- The *max-min-inequality* states that for any function $f(a, b)$

$$\max_a \min_b f(a, b) \leq \min_b \max_a f(a, b)$$

- Hence, the solution to the dual is a lower bound to the solution of the primal problem.
- The solution to the *dual function*, $\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu})$ is the min of a collection of linear functions, and thus always concave.
- It is easy to optimize this.
- In general, solving the dual is easy whenever minimizing \mathcal{L} wrt x is easy.

Penalty Methods

- We can convert the constrained problem back to unconstrained by adding penalty terms for constraint violations.
- A simple method could just count the number of violations:

$$p_{\text{count}}(\mathbf{x}) = \sum_i (h_i(\mathbf{x}) \neq 0) + \sum_j (g_j(\mathbf{x}) > 0)$$

- and add this to the objective in an *unconstrained* problem with penalty $\rho > 0$

$$\min_{\mathbf{x}} f(\mathbf{x}) + \rho p_{\text{count}}(\mathbf{x})$$

- One can choose the penalty function: for example, a quadratic penalty will produce a smooth objective function
- Notice that ρ needs to become very large sometimes here.

Augmented Lagrange Method

- This is very similar, but specific to equality constraints.

Interior Point Method

- Also called *barrier method*.
- These methods make sure that the search point remains always feasible.
- As one approaches the constraint boundary, the barrier function goes to infinity. Properties:
 1. $p_{\text{barrier}}(\mathbf{x})$ is continuous
 2. $p_{\text{barrier}}(\mathbf{x})$ is non negative
 3. $p_{\text{barrier}}(\mathbf{x})$ goes to infinity as one approaches the constraint boundary

2.5.2 Barriers

- Inverse Barrier

$$p_{\text{barrier}}(\mathbf{x}) = -\sum_i \frac{1}{g_i(\mathbf{x})}$$

- Log Barrier

$$p_{\text{barrier}}(\mathbf{x}) = -\sum_i \begin{cases} \log(-g_i(\mathbf{x})) & \text{if } g_i(\mathbf{x}) \geq -1 \\ 0 & \text{else.} \end{cases}$$

- The approach is as before, one transforms the problem to an unconstrained one and increases ρ until convergence:

$$\min_{\mathbf{x}} f(\mathbf{x}) + \frac{1}{\rho} p_{\text{barrier}}(\mathbf{x})$$

2.5.3 Examples

$$\min_{x \in \mathbb{R}^2} \sqrt{x_2} \text{ subject to } \begin{aligned} x_2 &\geq 0 \\ x_2 &\geq (a_1 x_1 + b_1)^3 \\ x_2 &\geq (a_2 x_1 + b_2)^3 \end{aligned}$$

2.6 Constrained Optimisation with `NLopt.jl`

- We need to specify one function for each objective and constraint.
- Both of those functions need to compute the function value (i.e. objective or constraint) *and* it's respective gradient.
- `NLopt` expects constraints **always** to be formulated in the format

$$g(x) \leq 0$$

where g is your constraint function

- The constraint function is formulated for each constraint at x . it returns a number (the value of the constraint at x), and it fills out the gradient vector, which is the partial derivative of the current constraint wrt x .
- There is also the option to have vector valued constraints, see the documentation.
- We set this up as follows:

In [27]: `using NLOpt`

```

count = 0 # keep track of # function evaluations

function myfunc(x::Vector, grad::Vector)
    if length(grad) > 0
        grad[1] = 0
        grad[2] = 0.5/sqrt(x[2])
    end

    global count
    count::Int += 1
    println("f_$(count)($x)")

    sqrt(x[2])
end

function myconstraint(x::Vector, grad::Vector, a, b)
    if length(grad) > 0
        grad[1] = 3a * (a*x[1] + b)^2
        grad[2] = -1
    end
    (a*x[1] + b)^3 - x[2]
end

opt = Opt{:LD_MMA, 2}
lower_bounds!(opt, [-Inf, 0.])
xtol_rel!(opt, 1e-4)

min_objective!(opt, myfunc)
inequality_constraint!(opt, (x,g) -> myconstraint(x,g,2,0), 1e-8)
inequality_constraint!(opt, (x,g) -> myconstraint(x,g,-1,1), 1e-8)

(minfunc,minx,ret) = NLOpt.optimize(opt, [1.234, 5.678])
println("got $minfunc at $minx after $count iterations (returned $ret)")

f_1([1.234, 5.678])
f_2([0.878739, 5.55137])
f_3([0.826216, 5.0439])
f_4([0.473944, 4.07677])
f_5([0.353898, 3.03085])
f_6([0.333873, 1.97179])

```

```

f_7([0.333334, 1.04509])
f_8([0.333334, 0.469503])
f_9([0.333333, 0.305792])
f_10([0.333333, 0.296322])
f_11([0.333333, 0.296296])
got 0.5443310477213124 at [0.333333, 0.296296] after 11 iterations (returned XTOL_REACHED)

```

WARNING: using NLOpt.optimize in module Main conflicts with an existing identifier.

2.7 NLOpt: Rosenbrock

- Let's tackle the rosenbrock example again.
- To make it more interesting, let's add an inequality constraint.

$$\min_{x \in \mathbb{R}^2} (1 - x_1)^2 + 100(x_2 - x_1^2)^2 \text{ subject to } 0.8 - x_1^2 - x_2^2 \geq 0$$

- in NLOpt format, the constraint is $x_1 + x_2 - 0.8 \leq 0$

```

In [28]: function rosenbrockf(x::Vector,grad::Vector)
           if length(grad) > 0
               grad[1] = -2.0 * (1.0 - x[1]) - 400.0 * (x[2] - x[1]^2) * x[1]
               grad[2] = 200.0 * (x[2] - x[1]^2)
           end
           return (1.0 - x[1])^2 + 100.0 * (x[2] - x[1]^2)^2
       end
       function r_constraint(x::Vector, grad::Vector)
           if length(grad) > 0
               grad[1] = 2*x[1]
               grad[2] = 2*x[2]
           end
           return x[1]^2 + x[2]^2 - 0.8
       end
       opt = Opt(:LD_MMA, 2)
       lower_bounds!(opt, [-5, -5.0])
       min_objective!(opt, (x,g) -> rosenbrockf(x,g))
       inequality_constraint!(opt, (x,g) -> r_constraint(x,g))
       ftol_rel!(opt, 1e-9)
       NLOpt.optimize(opt, [-1.0, 0.0])

```

```

Out [28]: (0.07588358473630112, [0.724702, 0.524221], :FTOL_REACHED)

```

2.8 JuMP.jl

- Introduce [JuMP.jl](#)
- JuMP is a mathematical programming interface for Julia. It is like AMPL, but for free and with a decent programming language.
- The main highlights are:

- It uses automatic differentiation to compute derivatives from your expression.
 - It supplies this information, as well as the sparsity structure of the Hessian to your preferred solver.
 - It decouples your problem completely from the type of solver you are using. This is great, since you don't have to worry about different solvers having different interfaces.
 - In order to achieve this, JuMP uses [MathProgBase.jl](#), which converts your problem formulation into a standard representation of an optimization problem.
- Let's look at the readme
 - The technical citation is Lubin et al [?]

2.9 JuMP: Quick start guide

- this is from the [quick start guide](#)
- please check the docs, they are excellent.

2.9.1 Step 1: create a model

- A model collects variables, objective function and constraints.
- it defines a solver to be used.

```
using Clp
m = Model(solver=ClpSolver()) # provide a solver

# Define variables
@variable(m, x ) # No bounds
@variable(m, x >= lb ) # Lower bound only (note: 'lb <= x' is not valid)
@variable(m, x <= ub ) # Upper bound only
@variable(m, lb <= x <= ub ) # Lower and upper bounds

# we can create arrays of a variable
N = 2
@variable(m, x[1:M,1:N] >= 0 )

# or put them in a block
@variables m begin
    x
    y >= 0
    Z[1:10], Bin
    X[1:3,1:3], SDP
    q[i=1:2], (lowerbound = i, start = 2i, upperbound = 3i)
    t[j=1:3], (Int, start = j)
end

# Equivalent to:
@variable(m, x)
@variable(m, y >= 0)
@variable(m, Z[1:10], Bin)
```

```

@variable(m, X[1:3,1:3], SDP)
@variable(m, q[i=1:2], lowerbound = i, start = 2i, upperbound = 3i)
@variable(m, t[j=1:3], Int, start = j)

# bounds can depend on indices
@variable(m, x[i=1:10] >= i )

```

2.10 Objective and Constraints

- We can easily add objective and constraint functions:

```

@constraint(m, x[i] - s[i] <= 0) # Other options: == and >=
@constraint(m, sum(x[i] for i=1:numLocation) == 1)
@objective(m, Max, 5x + 22y + (x+y)/2) # or Min

```

- This is fully integrated with Julia. you can use the generator syntax for sums:

```

@objective(sum(x[i] + y[i]/pi for i = I1, j = I2 if i+j < some_val))

```

```

In [29]: ##~Simple example
using JuMP
using Clp

let
    m = Model(solver = ClpSolver())
    @variable(m, 0 <= x <= 2 )
    @variable(m, 0 <= y <= 30 )

    @objective(m, Max, 5x + 3*y )
    @constraint(m, 1x + 5y <= 3.0 )

    print(m)

    status = solve(m)

    println("Objective value: ", getobjectivevalue(m))
    println("x = ", getvalue(x))
    println("y = ", getvalue(y))
end

```

```

Max 5 x + 3 y
Subject to
  x + 5 y  3
  0 x  2
  0 y  30
Objective value: 10.6
x = 2.0
y = 0.2

```



```

In [30]: # JuMP: Rosenbrock Example
         # Instead of hand-coding first and second derivatives, you only have to give `JuMP` e
         # Here is an example.

using Ipopt

let

    m = Model(solver=IpoptSolver())

    @variable(m, x)
    @variable(m, y)

    @NLOjective(m, Min, (1-x)^2 + 100(y-x^2)^2)

    solve(m)

    println("x = ", getvalue(x), " y = ", getvalue(y))

end

```

```

*****
This program contains Ipopt, a library for large-scale nonlinear optimization.
Ipopt is released as open source code under the Eclipse Public License (EPL).
For more information visit http://projects.coin-or.org/Ipopt
*****

```

This is Ipopt version 3.12.8, running with linear solver mumps.
NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

```

Number of nonzeros in equality constraint Jacobian...:      0
Number of nonzeros in inequality constraint Jacobian.:      0
Number of nonzeros in Lagrangian Hessian...:              3

```

```

Total number of variables...:      2
      variables with only lower bounds:      0
      variables with lower and upper bounds:  0
      variables with only upper bounds:      0
Total number of equality constraints...:      0
Total number of inequality constraints...:      0
      inequality constraints with only lower bounds:      0
      inequality constraints with lower and upper bounds:  0
      inequality constraints with only upper bounds:      0

```

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.0000000e+00	0.00e+00	2.00e+00	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	9.5312500e-01	0.00e+00	1.25e+01	-1.0	1.00e+00	-	1.00e+00	2.50e-01f	3

```

2  4.8320569e-01  0.00e+00  1.01e+00  -1.0  9.03e-02  -  1.00e+00  1.00e+00f  1
3  4.5708829e-01  0.00e+00  9.53e+00  -1.0  4.29e-01  -  1.00e+00  5.00e-01f  2
4  1.8894205e-01  0.00e+00  4.15e-01  -1.0  9.51e-02  -  1.00e+00  1.00e+00f  1
5  1.3918726e-01  0.00e+00  6.51e+00  -1.7  3.49e-01  -  1.00e+00  5.00e-01f  2
6  5.4940990e-02  0.00e+00  4.51e-01  -1.7  9.29e-02  -  1.00e+00  1.00e+00f  1
7  2.9144630e-02  0.00e+00  2.27e+00  -1.7  2.49e-01  -  1.00e+00  5.00e-01f  2
8  9.8586451e-03  0.00e+00  1.15e+00  -1.7  1.10e-01  -  1.00e+00  1.00e+00f  1
9  2.3237475e-03  0.00e+00  1.00e+00  -1.7  1.00e-01  -  1.00e+00  1.00e+00f  1
iter  objective    inf_pr  inf_du lg(mu)  ||d||  lg(rg) alpha_du alpha_pr  ls
10  2.3797236e-04  0.00e+00  2.19e-01  -1.7  5.09e-02  -  1.00e+00  1.00e+00f  1
11  4.9267371e-06  0.00e+00  5.95e-02  -1.7  2.53e-02  -  1.00e+00  1.00e+00f  1
12  2.8189505e-09  0.00e+00  8.31e-04  -2.5  3.20e-03  -  1.00e+00  1.00e+00f  1
13  1.0095040e-15  0.00e+00  8.68e-07  -5.7  9.78e-05  -  1.00e+00  1.00e+00f  1
14  1.3288608e-28  0.00e+00  2.02e-13  -8.6  4.65e-08  -  1.00e+00  1.00e+00f  1

```

Number of Iterations....: 14

```

                                (scaled)                (unscaled)
Objective....:  1.3288608467480825e-28  1.3288608467480825e-28
Dual infeasibility....:  2.0183854587685121e-13  2.0183854587685121e-13
Constraint violation....:  0.0000000000000000e+00  0.0000000000000000e+00
Complementarity....:  0.0000000000000000e+00  0.0000000000000000e+00
Overall NLP error....:  2.0183854587685121e-13  2.0183854587685121e-13

```

```

Number of objective function evaluations      = 36
Number of objective gradient evaluations     = 15
Number of equality constraint evaluations     = 0
Number of inequality constraint evaluations   = 0
Number of equality constraint Jacobian evaluations = 0
Number of inequality constraint Jacobian evaluations = 0
Number of Lagrangian Hessian evaluations    = 14
Total CPU secs in IPOPT (w/o function evaluations) = 0.158
Total CPU secs in NLP function evaluations    = 0.024

```

EXIT: Optimal Solution Found.

x = 0.99999999999999899 y = 0.9999999999999792

```

In [31]: # not bad, right?
         # adding the constraint from before:

```

```
let
```

```
    m = Model(solver=IpoptSolver())
```

```
    @variable(m, x)
```

```
    @variable(m, y)
```

```

@NLObjective(m, Min, (1-x)^2 + 100(y-x^2)^2)
@NLconstraint(m,x^2 + y^2 <= 0.8)

solve(m)

println("x = ", getvalue(x), " y = ", getvalue(y))

end

```

This is Ipopt version 3.12.8, running with linear solver mumps.
NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

```

Number of nonzeros in equality constraint Jacobian...:      0
Number of nonzeros in inequality constraint Jacobian.:      2
Number of nonzeros in Lagrangian Hessian...:              5

```

```

Total number of variables...:      2
      variables with only lower bounds:      0
      variables with lower and upper bounds:    0
      variables with only upper bounds:      0
Total number of equality constraints...:      0
Total number of inequality constraints...:      1
      inequality constraints with only lower bounds:      0
      inequality constraints with lower and upper bounds:    0
      inequality constraints with only upper bounds:      1

```

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.0000000e+00	0.00e+00	2.00e+00	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	9.5312500e-01	0.00e+00	1.25e+01	-1.0	5.00e-01	-	1.00e+00	5.00e-01f	2
2	4.9204994e-01	0.00e+00	9.72e-01	-1.0	8.71e-02	-	1.00e+00	1.00e+00f	1
3	2.0451702e+00	0.00e+00	3.69e+01	-1.7	3.80e-01	-	1.00e+00	1.00e+00H	1
4	1.0409466e-01	0.00e+00	3.10e-01	-1.7	1.46e-01	-	1.00e+00	1.00e+00f	1
5	8.5804626e-02	0.00e+00	2.71e-01	-1.7	9.98e-02	-	1.00e+00	1.00e+00h	1
6	9.4244879e-02	0.00e+00	6.24e-02	-1.7	3.74e-02	-	1.00e+00	1.00e+00h	1
7	8.0582034e-02	0.00e+00	1.51e-01	-2.5	6.41e-02	-	1.00e+00	1.00e+00h	1
8	7.8681242e-02	0.00e+00	2.12e-03	-2.5	1.12e-02	-	1.00e+00	1.00e+00h	1
9	7.6095770e-02	0.00e+00	6.16e-03	-3.8	1.37e-02	-	1.00e+00	1.00e+00h	1
iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
10	7.6033892e-02	0.00e+00	2.23e-06	-3.8	3.99e-04	-	1.00e+00	1.00e+00h	1
11	7.5885642e-02	0.00e+00	2.07e-05	-5.7	7.99e-04	-	1.00e+00	1.00e+00h	1
12	7.5885428e-02	0.00e+00	2.74e-11	-5.7	1.38e-06	-	1.00e+00	1.00e+00h	1
13	7.5883585e-02	0.00e+00	3.19e-09	-8.6	9.93e-06	-	1.00e+00	1.00e+00f	1

Number of Iterations...: 13

```

                                (scaled)                (unscaled)
Objective...:  7.5883585442440671e-02  7.5883585442440671e-02

```

```

Dual infeasibility...: 3.1949178858070582e-09    3.1949178858070582e-09
Constraint violation...: 0.0000000000000000e+00    0.0000000000000000e+00
Complementarity...: 2.5454985882932001e-09    2.5454985882932001e-09
Overall NLP error...: 3.1949178858070582e-09    3.1949178858070582e-09

```

```

Number of objective function evaluations      = 20
Number of objective gradient evaluations     = 14
Number of equality constraint evaluations     = 0
Number of inequality constraint evaluations   = 20
Number of equality constraint Jacobian evaluations = 0
Number of inequality constraint Jacobian evaluations = 14
Number of Lagrangian Hessian evaluations    = 13
Total CPU secs in IPOPT (w/o function evaluations) = 0.005
Total CPU secs in NLP function evaluations   = 0.003

```

EXIT: Optimal Solution Found.

x = 0.7247018392092258 y = 0.5242206029480763

2.11 JuMP: Maximum Likelihood

- Let's redo the maximum likelihood example in JuMP.
- Let μ, σ^2 be the unknown mean and variance of a random sample generated from the normal distribution.
- Find the maximum likelihood estimator for those parameters!
- density:

$$f(x_i|\mu, \sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x_i - \mu)^2}{2\sigma^2}\right)$$

- Likelihood Function

$$\begin{aligned} L(\mu, \sigma^2) &= \prod_{i=1}^N f(x_i|\mu, \sigma^2) = \frac{1}{(\sigma\sqrt{2\pi})^n} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2\right) \\ &= (\sigma^2 2\pi)^{-\frac{n}{2}} \exp\left(-\frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2\right) \end{aligned}$$

- Constraints: $\mu \in \mathbb{R}, \sigma > 0$
- log-likelihood:

$$\log L = l = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^N (x_i - \mu)^2$$

- Let's do this in JuMP.

```

In [32]: # Copyright 2015, Iain Dunning, Joey Huchette, Miles Lubin, and contributors
# example modified
using Distributions

let
    distrib = Normal(4.5,3.5)
    n = 10000

    data = rand(distrib,n);

    m = Model(solver=IpoptSolver())

    @variable(m, mu, start = 0.0)
    @variable(m, sigma >= 0.0, start = 1.0)

    @NLobjective(m, Max, -(n/2)*log(2*sigma^2)-sum((data[i] - mu) ^ 2 for i = 1:n)/(2))

    solve(m)
    println(" = ", getvalue(mu),",", mean(data) = ", mean(data))
    println("^2 = ", getvalue(sigma)^2, ",", var(data) = ", var(data))
end

```

This is Ipopt version 3.12.8, running with linear solver mumps.

NOTE: Other linear solvers might be more efficient (see Ipopt documentation).

```

Number of nonzeros in equality constraint Jacobian...:      0
Number of nonzeros in inequality constraint Jacobian.:      0
Number of nonzeros in Lagrangian Hessian...:           3

```

```

Total number of variables...:          2
      variables with only lower bounds:      1
      variables with lower and upper bounds:  0
      variables with only upper bounds:      0
Total number of equality constraints...:    0
Total number of inequality constraints...:  0
      inequality constraints with only lower bounds:      0
      inequality constraints with lower and upper bounds:  0
      inequality constraints with only upper bounds:      0

```

iter	objective	inf_pr	inf_du	lg(mu)	d	lg(rg)	alpha_du	alpha_pr	ls
0	1.7210445e+05	0.00e+00	1.01e+02	-1.0	0.00e+00	-	0.00e+00	0.00e+00	0
1	1.2939451e+05	0.00e+00	1.07e+02	-1.0	9.80e+00	-	1.00e+00	5.00e-01f	2
2	7.8864616e+04	0.00e+00	4.43e+01	-1.0	2.27e-01	-	8.88e-01	1.00e+00f	1
3	5.1720095e+04	0.00e+00	1.84e+01	-1.0	2.96e-01	-	1.00e+00	1.00e+00f	1
4	3.7648943e+04	0.00e+00	7.44e+00	-1.0	3.75e-01	-	1.00e+00	1.00e+00f	1
5	3.0796009e+04	0.00e+00	2.93e+00	-1.0	4.57e-01	-	1.00e+00	1.00e+00f	1
6	2.7858621e+04	0.00e+00	1.07e+00	-1.0	5.13e-01	-	1.00e+00	1.00e+00f	1
7	2.6886136e+04	0.00e+00	3.39e-01	-1.0	4.94e-01	-	1.00e+00	1.00e+00f	1

```

      8  2.6711082e+04  0.00e+00  6.36e-02  -1.7  3.19e-01    -  1.00e+00  1.00e+00f  1
      9  2.6701110e+04  0.00e+00  4.25e-03  -2.5  1.01e-01    -  1.00e+00  1.00e+00f  1
iter   objective      inf_pr   inf_du lg(mu)  ||d||  lg(rg) alpha_du alpha_pr ls
     10  2.6701072e+04  0.00e+00  1.72e-05  -3.8  6.78e-03    -  1.00e+00  1.00e+00f  1
     11  2.6701072e+04  0.00e+00  8.12e-10  -5.7  4.62e-05    -  1.00e+00  1.00e+00f  1
     12  2.6701072e+04  0.00e+00  4.04e-13  -8.6  1.02e-06    -  1.00e+00  1.00e+00f  1

```

Number of Iterations....: 12

```

                                (scaled)                (unscaled)
Objective....:  8.4542511167531504e+00  2.6701072346234618e+04
Dual infeasibility....:  4.0444162230182163e-13  1.2773485040573046e-09
Constraint violation....:  0.0000000000000000e+00  0.0000000000000000e+00
Complementarity....:  2.5064398796454131e-09  7.9160923461665668e-06
Overall NLP error....:  2.5064398796454131e-09  7.9160923461665668e-06

```

```

Number of objective function evaluations      = 18
Number of objective gradient evaluations      = 13
Number of equality constraint evaluations      = 0
Number of inequality constraint evaluations    = 0
Number of equality constraint Jacobian evaluations = 0
Number of inequality constraint Jacobian evaluations = 0
Number of Lagrangian Hessian evaluations      = 12
Total CPU secs in IPOPT (w/o function evaluations) = 0.007
Total CPU secs in NLP function evaluations      = 0.026

```

EXIT: Optimal Solution Found.

```

= 4.513536352090578, mean(data) = 4.513536352090578
^2 = 12.211002671496686, var(data) = 12.212223884224223

```

3 Linear Constrained Problems (LPs)

- Very similar to before, just that both objective and constraints are *linear*.

$$\begin{aligned}
 & \min_{\mathbf{x}} \mathbf{c}^T \mathbf{x} \\
 & \text{subject to } \mathbf{w}_{LE}^{(i)T} \mathbf{x} \leq b_i \text{ for } i \in 1, 2, 3, \dots \\
 & \quad \mathbf{w}_{GE}^{(j)T} \mathbf{x} \geq b_j \text{ for } j \in 1, 2, 3, \dots \\
 & \quad \mathbf{w}_{EQ}^{(k)T} \mathbf{x} = b_k \text{ for } k \in 1, 2, 3, \dots
 \end{aligned}$$

3.0.1 Standard Form

- Usually LPs are given in *standard form*
- All constraints are less-than inequalities

- All choice variables are non-negative.

$$\begin{aligned} \min_{\mathbf{x}} \mathbf{c}^T \mathbf{x} \\ \text{subject to } \mathbf{Ax} \leq \mathbf{b} \\ \mathbf{x} \geq 0 \end{aligned}$$

- Greater-than inequality constraints are inverted
- equality constraints are split into two
- $\mathbf{x} = \mathbf{x}^+ - \mathbf{x}^-$ and we constrain both components to be positive.

3.0.2 Equality Form

$$\begin{aligned} \min_{\mathbf{x}} \mathbf{c}^T \mathbf{x} \\ \text{subject to } \mathbf{Ax} = \mathbf{b} \\ \mathbf{x} \geq 0 \end{aligned}$$

- Can transform standard into equality form

$$\mathbf{Ax} \leq \mathbf{b} \rightarrow \mathbf{Ax} + \mathbf{s} = \mathbf{b}, \mathbf{s} \geq 0$$

- equality constraints are split into two
- $\mathbf{x} = \mathbf{x}^+ - \mathbf{x}^-$ and we constrain both components to be positive.

3.0.3 Solving LPs

- Simplex Algorithm operates on Equality Form
- Moving from one vertex to the next of the feasible set, this is guaranteed to find the optimal solution if the problem is bounded.

4 Discrete Optimization / Integer Programming

- Here the choice variable is constrained to come from a discrete set \mathcal{X} .
- If this set is $\mathcal{X} = \mathbb{N}$, we have an **integer program**
- If only *some* x have to be discrete, this is a **mixed integer program**

4.1 Example

$$\begin{aligned} \min_{\mathbf{x}} x_1 + x_2 \\ \text{subject to } \|\mathbf{x}\| \leq 2 \\ \mathbf{x} \in \mathbb{N} \end{aligned}$$

- continuous optimum is $(-\sqrt{2}, -\sqrt{2})$ and objective is $y = -2\sqrt{2} = -2.828$
- Integer constrained problem is only delivering $y = -2$, and $\mathbf{x}^* \in (-2, 0), (-1, -1), (0, -2)$

```
In [33]: x = -3:0.01:3
dx = repmat(linspace(-3,3,7),1,7)
contourf(x,x,(x,y)->x+y,color=:blues)
scatter!(dx,dx',legend=false,markercolor=:white)
plot!(x->sqrt(4-x^2),-2,2,c=:white)
plot!(x->-sqrt(4-x^2),-2,2,c=:white)
scatter!([-2,-1,0],[0,-1,-2],c=:red)
scatter!([-sqrt(2)],[-sqrt(2)],c=:red,markershape=:cross,markersize=9)
```

4.2 Rounding

- One solution is to just *round the continuous solution to the nearest integer*
- We compute the **relaxed** problem, i.e. the one where x is continuous.
- Then we round up or down.
- Can go terribly wrong.

4.3 Cutting Planes

- This is an exact method
- We solve the relaxed problem first.
- Then we add linear constraints that result in the solution becoming integral.

4.4 Branch and Bound

- This enumerates all possible solutions.
- Branch and bound does this, without having to compute all of them.

4.5 Example: The Knapsack Problem

- We are packing our knapsack for a trip but only have space for the most valuable items.
- We have $x_i = 0$ if item i is not in the sack, 1 else.

$$\begin{aligned} \min_x \quad & \sum_{i=1}^n v_i x_i \\ \text{s.t.} \quad & \sum_{i=1}^n w_i x_i \leq w_{\max} \\ & w_i \in \mathbb{N}_+, v_i \in \mathbb{R} \end{aligned}$$

- If there are n items, we have 2^n possible design vectors.
- But there is a useful recursive relationship.
- If we solved $n - 1$ knapsack problems so far and deliberate about item n
 - If it's not worth including item n , then the solution is the knapsack problem for $n - 1$ items and capacity w_{\max}
 - If it IS worth including it: solution will have value of knapsack with $n - 1$ items and reduced capacity, plus the value of the new item


```

In [34]: # Copyright 2017, Iain Dunning, Joey Huchette, Miles Lubin, and contributors
# This Source Code Form is subject to the terms of the Mozilla Public
# License, v. 2.0. If a copy of the MPL was not distributed with this
# file, You can obtain one at http://mozilla.org/MPL/2.0/.
#####
# JuMP
# An algebraic modeling language for Julia
# See http://github.com/JuliaOpt/JuMP.jl
#####
# knapsack.jl
#
# Solves a simple knapsack problem:
# max sum(p_j x_j)
# st sum(w_j x_j) <= C
# x binary
#####

using JuMP, Cbc

let

    # Maximization problem
    m = Model(solver=CbcSolver())

    @variable(m, x[1:5], Bin)

    profit = [ 5, 3, 2, 7, 4 ]
    weight = [ 2, 8, 4, 2, 5 ]
    capacity = 10

    # Objective: maximize profit
    @objective(m, Max, dot(profit, x))

    # Constraint: can carry all
    @constraint(m, dot(weight, x) <= capacity)

    # Solve problem using MIP solver
    status = solve(m)

    println("Objective is: ", getobjectivevalue(m))
    println("Solution is:")
    for i = 1:5
        print("x[$i] = ", getvalue(x[i]))
        println(", p[$i]/w[$i] = ", profit[i]/weight[i])
    end
end
end

```

Objective is: 16.0

Solution is:

```
x[1] = 1.0, p[1]/w[1] = 2.5  
x[2] = 0.0, p[2]/w[2] = 0.375  
x[3] = 0.0, p[3]/w[3] = 0.5  
x[4] = 1.0, p[4]/w[4] = 3.5  
x[5] = 1.0, p[5]/w[5] = 0.8
```

In []: